

# Process Environment

---

Advanced Programming Environment in the UNIX Environment

Chun-Ying Huang <chuang@cs.nctu.edu.tw>

# Outline

---

Process start and termination

Environment variables

Memory layout

Shared libraries

Memory allocation

setjmp and longjmp

Process resource limits

# Process Start

---

The main function

Synopsis

- `int main(int argc, char *argv[]);`
- `int main(int argc, char *argv[], char *envp[]);`

# Process Termination

---

## Normal process termination in five ways

- Return from main
- Calling exit
- Calling \_exit or \_Exit
- Return of the last thread from its start routine
- Calling pthread\_exit from the last thread

## Abnormal process termination in three ways

- Calling abort
- Receipt of a signal
- Response of the last thread to a cancellation request

## Execution of a main function looks like

- `exit(main(argc, argv));`

# atexit and exit Functions

---

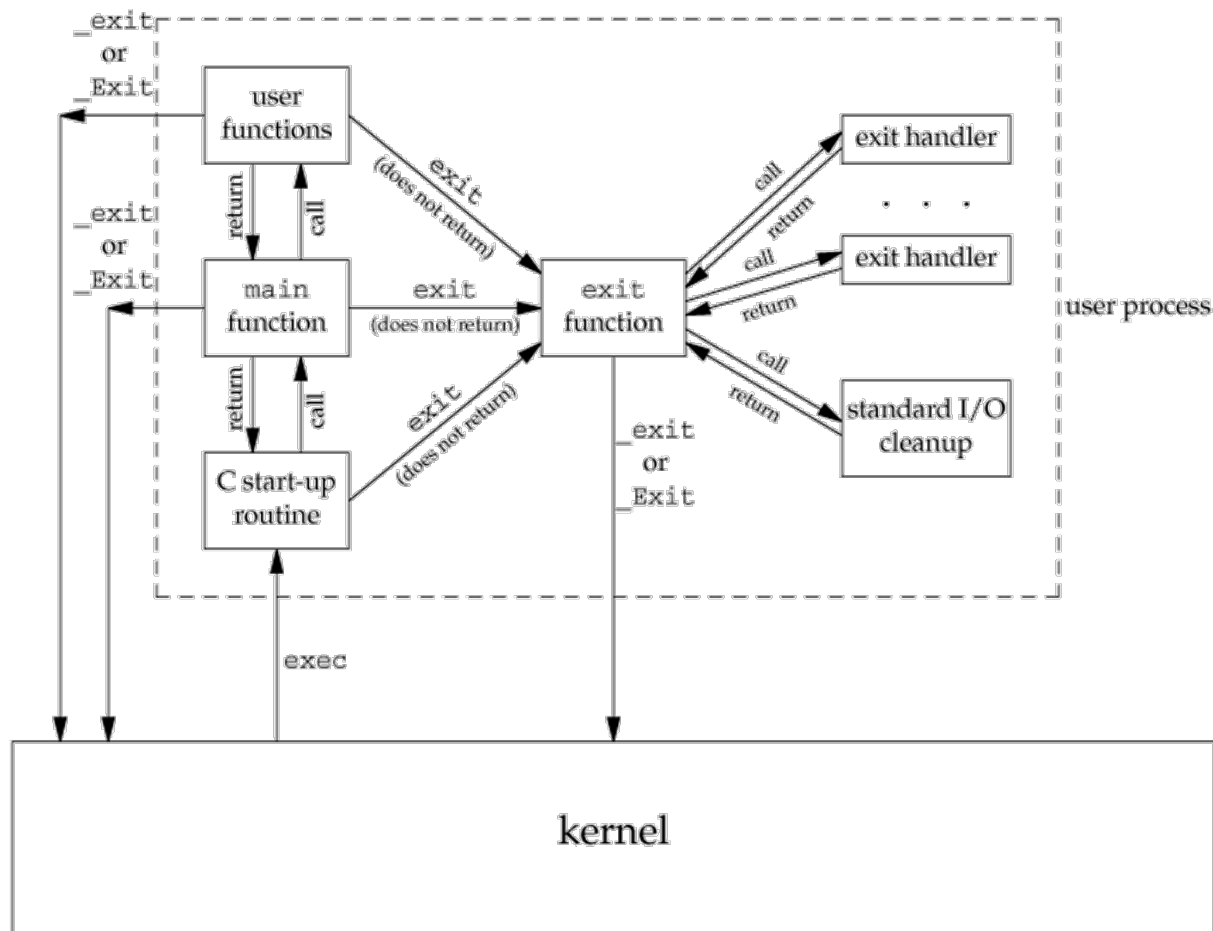
## Manual cleanups on exit

- `int atexit(void (*function)(void));`
- Register up to 32 customized functions (textbook)
  - Linux has extended this restrictions

## Exit functions

- `exit`
  - Call atexit registered functions
  - Performed a clean shutdown of the standard I/O library
  - `fclose()` all streams, `remove tmpfile()`
- `_exit` and `_Exit`
  - Terminate immediately

# Start and Termination of a C Program



# Environment Variables

---

## The environment variables

- Usually in the form of: `name=value` (no spaces around =)
- Relevant commands: `env`, `export` (bash)
- Use `$` to read a specific environment variable in a shell

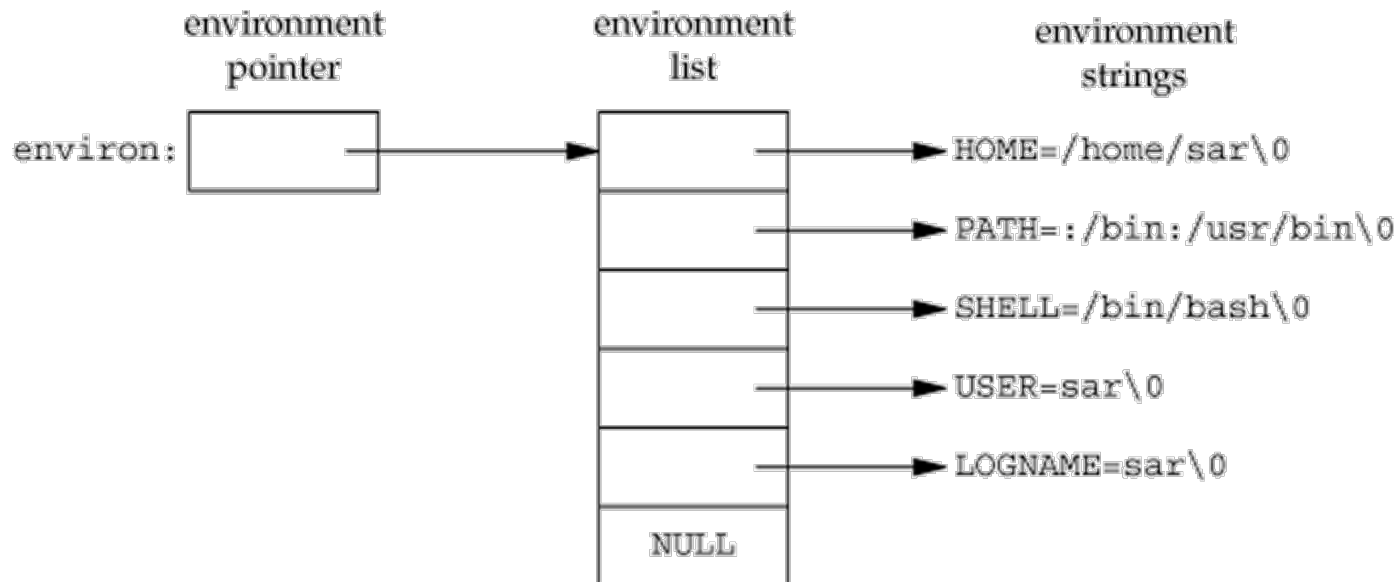
## List of environment variable functions

Function	ISO C	POSIX.1	FreeBSD 8.0	Linux 3.2.0	Mac OS X 10.6.8	Solaris 10
<code>getenv</code>	•	•	•	•	•	•
<code>putenv</code>		XSI	•	•	•	•
<code>setenv</code>		•	•	•	•	
<code>unsetenv</code>		•	•	•	•	
<code>clearenv</code>				•		

# Environment List

Access environment variables directly

- `int main(int argc, char *argv[], char *envp[]);`
- `extern char **environ;`





# Environment Functions

---

Prototypes of functions to manipulate environment variables

```
#include <stdlib.h>
```

```
char *getenv(const char *name);
```

```
int putenv(char *string);
```

```
int setenv(const char *name, const char *value, int overwrite);
```

```
int unsetenv(const char *name);
```

```
int clearenv(void);
```

# Environment List Operations

---

## Delete an entry

- This is simple, just free a string and move all subsequent pointers down one

## Modify an entry

- If new-size  $\geq$  old-size, just overwrite the old one
- If new-size  $>$  old-size, allocate a new space the new variable and make the pointer point to the new location

## Add an entry

- Add for the 1<sup>st</sup> time, allocate a new space for the entire list
- Add for non-1<sup>st</sup> time, reallocate a larger space for the entire list

# Common Environment Variables (1/3)

---

Variable	POSIX.1	FreeBSD 8.0	Linux 3.2.0	Mac OS X 10.6.8	Solaris 10	Description
COLUMNS	•	•	•	•	•	Terminal width
DATETIME	XSI		•	•	•	getdate(3) template file pathname
HOME	•	•	•	•	•	Home directory
LANG	•	•	•	•	•	Name of locale
LC_ALL	•	•	•	•	•	Name of locale
LC_COLLATE	•	•	•	•	•	Name of locale for collation
LC_CTYPE	•	•	•	•	•	Name of locale for character classification
LC_MESSAGES	•	•	•	•	•	Name of locale for messages

# Common Environment Variables (2/3)

---

Variable	POSIX.1	FreeBSD 8.0	Linux 3.2.0	Mac OS X 10.6.8	Solaris 10	Description
LC_MONETARY	•	•	•	•	•	Name of locale for monetary editing
LC_NUMERIC	•	•	•	•	•	Name of locale for numeric editing
LC_TIME	•	•	•	•	•	Name of locale for date/time formatting
LINES	•	•	•	•	•	Terminal height
LOGNAME	•	•	•	•	•	Login name
MSGVERB	XSI	•	•	•	•	fmtmsg(3) message components to process
NLSPATH	•	•	•	•	•	Sequence of templates for message catalogs

# Common Environment Variables (3/3)

---

Variable	POSIX.1	FreeBSD 8.0	Linux 3.2.0	Mac OS X 10.6.8	Solaris 10	Description
PATH	•	•	•	•	•	List of path prefixes to search for executable file
PWD	•	•	•	•	•	Absolute pathname of current working directory
SHELL	•	•	•	•	•	Name of user's preferred shell
TERM	•	•	•	•	•	Terminal type
TMPDIR	•	•	•	•	•	Pathname of directory for creating temporary files
TZ	•	•	•	•	•	Time zone information

# Memory Layout of a Program

## Text segment

- Machine instructions

## Initialized data segment

- `int maxcount = 100;`

## Uninitialized data segment (bss)

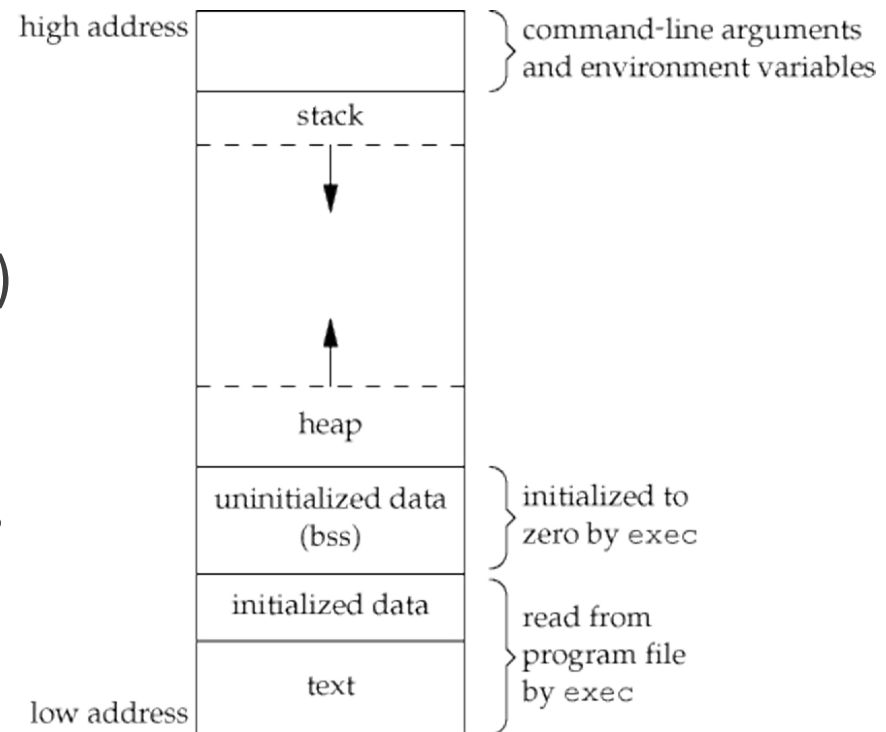
- `long sum[1000];`

## Stack

- Local variables, function call states

## Heap

- Dynamic allocated memory



# Guess the Number

```
huangant@rhino: ~/myprog/unix_prog/procenv
7 int main() {
8     char buf[16];
9     int answer;
10
11     setvbuf(stdout, NULL, _IONBF, 0);
12     srand(time(0) ^ getpid());
13     answer = rand() % 10000;
14
15     printf("Guess the number: ");
16     if(fgets(buf, 32, stdin) != NULL) {
17         int g = strtol(buf, NULL, 0);
18         printf("Your guess is %d\n", g);
19         if(g == answer) {
20             printf("Bingo!\n");
21         } else {
22             printf("No no no ... \n");
23         }
24     }
25
26     return 0;
27 }
```

NORMAL >> guess.c[+] c << 96% ln :27/28≡%:2 ≡ [27]tra...

guess.c: compile it with **-fno-stack-protector** option  
runtime: python3-pip, python3-virtualenv, and pwntools

# Read Sizes of an Executable Binary

---

The size (1) command

```
$ size /usr/bin/gcc /bin/sh
```

text	data	bss	dec	hex	filename
203913	2152	2248	208313	32db9	/usr/bin/gcc
704028	19268	19736	743032	b5678	/bin/sh



# Shared Libraries

---

Most UNIX systems today support shared libraries

Shared libraries remove the common library routines from the executable file

Maintain a single copy of the library routine somewhere in memory that all processes reference

- Reduce the size and memory requirement of each executable file
- But It may add some runtime overhead

Another advantage of shared libraries

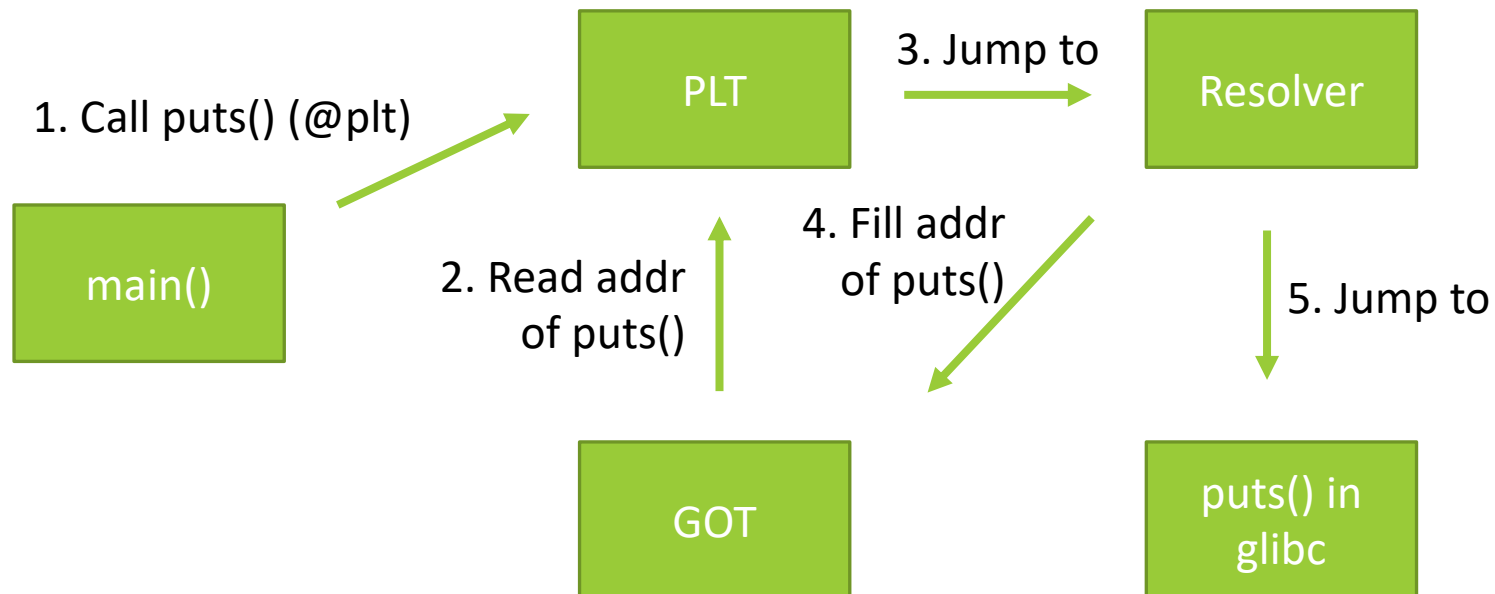
- Library functions can be replaced with new versions without having to relink every program that uses the library
- But it might also be a security flaw

# Resolve Functions in Shared Libraries: call puts() – 1st Time

GOT – global offset table: Store the "real address" of a function

PLT – procedure linkage table: Call the real address, or resolve it!

Resolver – Procedures to resolve the "real address" of a function

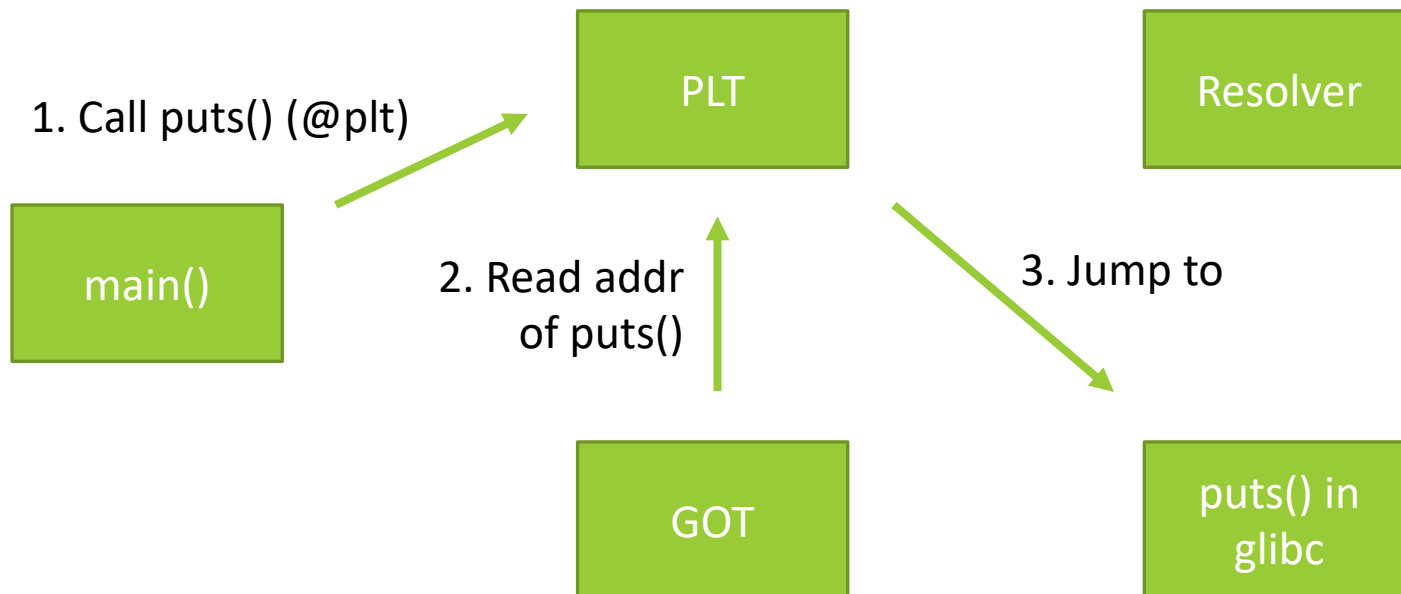


# Resolve Functions in Shared Libraries: call puts() – 2nd Time +

GOT – global offset table: Store the "real address" of a function

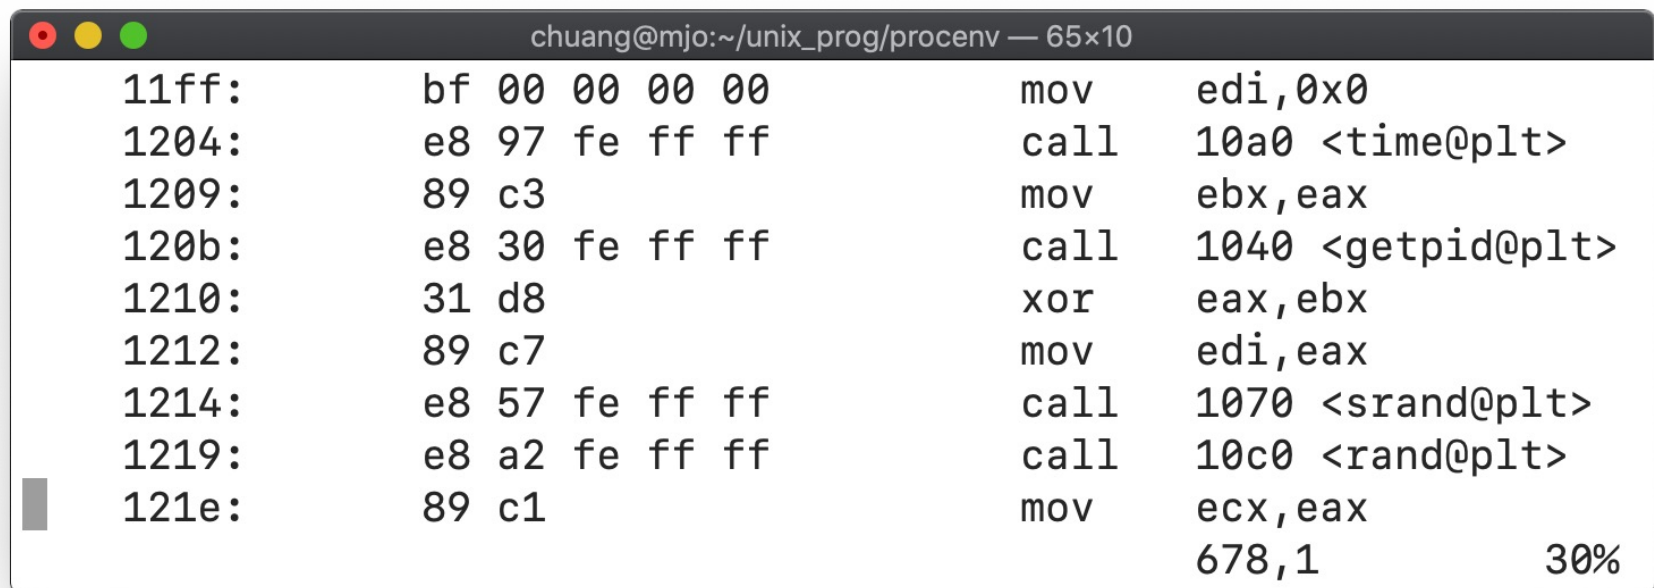
PLT – procedure linkage table: Call the real address, or resolve it!

Resolver – Procedures to resolve the "real address" of a function



# Resolving Functions in Shared Libraries – Sample: call time()

```
14      srand(time(0) ^ getpid());
```



The screenshot shows a debugger window titled 'chuang@mjo:~/unix\_prog/procenv — 65x10'. The window displays assembly code for a function, likely 'call time()'. The code is organized into columns: address, hex bytes, and assembly instructions. The instructions include 'mov edi, 0x0', 'call 10a0 <time@plt>', 'mov ebx, eax', 'call 1040 <getpid@plt>', 'xor eax, ebx', 'mov edi, eax', 'call 1070 <srand@plt>', 'call 10c0 <rand@plt>', and 'mov ecx, eax'. The address 121e: is highlighted. The instruction 'mov ecx, eax' is followed by the value '678,1' and a percentage '30%'.

Address	Hex Bytes	Assembly Instruction
11ff:	bf 00 00 00 00	mov edi, 0x0
1204:	e8 97 fe ff ff	call 10a0 <time@plt>
1209:	89 c3	mov ebx, eax
120b:	e8 30 fe ff ff	call 1040 <getpid@plt>
1210:	31 d8	xor eax, ebx
1212:	89 c7	mov edi, eax
1214:	e8 57 fe ff ff	call 1070 <srand@plt>
1219:	e8 a2 fe ff ff	call 10c0 <rand@plt>
121e:	89 c1	mov ecx, eax

678,1 30%

From guess.c/guess.s, addresses might be different from your build

# Resolving Functions in Shared Libraries – Sample: GOT & PLT

```
chuang@mjo:~/unix_prog/procenv — 103x11
00000000000010a0 <time@plt>:
  10a0:    ff 25 aa 2f 00 00      jmp     QWORD PTR [rip+0x2faa]      # 4050 <time@GLIBC_2.2.5>
  10a6:    68 07 00 00 00      push    0x7
  10ab:    e9 70 ff ff ff      jmp     1020 <.plt>

00000000000010c0 <rand@plt>:
  10c0:    ff 25 9a 2f 00 00      jmp     QWORD PTR [rip+0x2f9a]      # 4060 <rand@GLIBC_2.2.5>
  10c6:    68 09 00 00 00      push    0x9
  10cb:    e9 50 ff ff ff      jmp     1020 <.plt>
```

**PLT**  
574,0-1 25%

```
chuang@mjo:~/unix_prog/procenv — 103x13
4049:    10 00      adc     BYTE PTR [rax],al
404b:    00 00      add     BYTE PTR [rax],al
404d:    00 00      add     BYTE PTR [rax],al
404f:    00 a6 10 00 00 00      add     BYTE PTR [rsi+0x10],ah
4055:    00 00      add     BYTE PTR [rax],al
4057:    00 b6 10 00 00 00      add     QWORD PTR [rsi+0x10],dh
405d:    00 00      add     QWORD PTR [rax],al
405f:    00 c6      add     QWORD PTR [rax],al
4061:    10 00      adc     BYTE PTR [rax],al
4063:    00 00      add     BYTE PTR [rax],al
4065:    00 00      add     BYTE PTR [rax],al
...
```

**GOT**  
1174,1-8 52%

# Resolving Functions in Shared Libraries – Sample: DEMO

---

## Relevant GDB commands

### Display information

- `x/30i` – dump assembly instructions, e.g., `x/30i main`
- `x/gx` – dump a single 64-bit value in hexadecimal format
- `tui enable` – enable gdb's text user interface, "Ctrl-x o" to switch focus
- `layout asm` – change TUI layout to display assembly codes
- `layout src` – change TUI layout to display source codes

### Control flow

- `b` – set break points, e.g, `b *main+59`
- `run` – run the program
- `start` – run the first line of the program, and then paused
- `starti` – run the first instruction of the program, and then paused
- `si` – step one instruction
- `ni` – step one instruction, but proceed through subroutine calls
- `watch` – monitor memory writing

# Compile Static and Dynamic Program

---

A simple program that just print “Hello, World!”

```
$ gcc h1.c -o h1
$ gcc h2.c -o h2 -static
$ ls -la h1 h2
-rwxrwxr-x 1  chuang  chuang    9564 Mar 13 11:48 h1
-rwxrwxr-x 1  chuang  chuang 878192 Mar 13 11:48 h2
$ size h1 h2
   text    data     bss      dec       hex filename
   896     264        8    1168     490 h1
499650   1928    6948  508526   7c26e h2
```

# Library Injection

---

Functions referenced to shared libraries can be overridden

- The LD\_PRELOAD environment variable
- Usage:  
LD\_PRELOAD=/path/to/the/injected-shared-object {program}

Library injection does not work for suid/sgid executables



# Library Injection Example

---

Suppose we are going to hijack the `getuid()` function

- This is commonly used in tools like fake-root

The original program (`getuid.c`)

```
int main() {  
    printf("UID = %d\n", getuid());  
    return 0;  
}
```

The injected library (`inject1.c`)

```
#include <stdio.h>  
#include <sys/types.h>  
  
uid_t getuid(void) {  
    fprintf(stderr, "injected getuid, always return 0\n");  
    return 0;  
}
```

# Library Injection Example (Cont'd)

---

Compile the programs and the libraries

```
$ gcc -o getuid -Wall -g getuid.c
$ gcc -o inject1.so -shared -fPIC inject1.c -ldl
```

- The first command produces the getuid program
- The second commands generates the inject1.so (shared) library

Run the example

```
$ ./getuid                                # no injection
UID = 1000
$ LD_PRELOAD=./inject1.so ./getuid        # injected
injected getuid, always return 0
UID = 0
```

# More on Library Injection

---

But we still want the original function to work properly

We have to locate the original function

```
#include <dlfcn.h>

void *dlopen(const char *filename, int flag);
char *dlerror(void);
void *dlsym(void *handle, const char *symbol);
int dlclose(void *handle);
```

You may have to link with `-ldl` option

# Revised Library Injection Example

---

We would like to know the real UID internally (inject2.c)

```
#include <dlfcn.h>
#include <stdio.h>
#include <sys/types.h>

static uid_t (*old_getuid)(void) = NULL;    /* function pointer */

uid_t getuid(void) {
    if(old_getuid == NULL) {
        void *handle = dlopen("libc.so.6", RTLD_LAZY);
        if(handle != NULL)
            old_getuid = dlsym(handle, "getuid");
    }
    fprintf(stderr, "injected getuid, always return 0\n");
    if(old_getuid != NULL)
        fprintf(stderr, "real uid = %d\n", old_getuid());
    return 0;
}
```

# Revised Library Injection Example (Cont'd)

---

Compile the programs and the libraries (again)

```
$ gcc -o getuid -Wall -g getuid.c
$ gcc -o inject2.so -shared -fPIC inject2.c -ldl
```

- The first command produces the getuid program
- The second commands generates the inject2.so (shared) library

Run the example

```
$ ./getuid                                # no injection
UID = 1000
$ LD_PRELOAD=./inject2.so ./getuid        # injected
injected getuid, always return 0
real uid = 1000
UID = 0
```

# Determine Library Injection Possibility

---

No SUID/SGID enabled

Not a statically linked binary

Examples of the dynamic/static linked hello-world example

- The file command

```
$ file h1 h2
```

```
h1: ELF 64-bit LSB executable, x86-64, version 1 (SYSV), dynamically linked (uses shared  
libs), for GNU/Linux 2.6.24, BuildID[sha1]=e32f08cfbdda94d57273829c2bfd535d8fbe626d, not  
stripped
```

```
h2: ELF 64-bit LSB executable, x86-64, version 1 (GNU/Linux), statically linked, for  
GNU/Linux 2.6.24, BuildID[sha1]=2748d80822e76d183d0ef5633c0b784527727c7a, not stripped
```

- The ldd command

```
$ ldd h1 h2
```

```
h1:
```

```
linux-vdso.so.1 => (0x00007ffe7d3d5000)  
libc.so.6 => /lib/x86_64-linux-gnu/libc.so.6 (0x00007f1bc2150000)  
/lib64/ld-linux-x86-64.so.2 (0x00007f1bc2515000)
```

```
h2:
```

```
not a dynamic executable
```

# Determine Library Injection Possibility (Cont'd)

---

Use symbols from a shared library

The `nm` command

Example: static VS dynamic linked symbols

```
$ gcc -o getuid -Wall -g getuid.c                # dynamically linked
$ gcc -o getuid_s -Wall -g getuid.c -static      # statically linked
$ nm getuid | grep getuid
                 U getuid@@GLIBC_2.2.5          # getuid is unknown
$ nm getuid_s | grep getuid
0000000000433590 W getuid                        # getuid is known (but weak)
0000000000433590 T __getuid                     # the getuid implementation
```

Symbols can be stripped using the `strip` command

# Memory Allocation

---

## ISO C memory allocation functions

`void *malloc(size_t size);`

- Allocates a specified number of bytes of memory
- The initial value of the memory is indeterminate

`void *calloc(size_t nobj, size_t size);`

- Allocates space for a specified number of objects of a specified size
- The space is initialized to all 0 bits

`void *realloc(void *ptr, size_t newsize);`

- Increases or decreases the size of a previously allocated area
- It may involve moving the previously allocated area somewhere else, to provide the additional room at the end
- The initial value of increased memory is indeterminate



# Memory Allocation (Cont'd)

---

Allocated memory can be released by `free()`

The allocation routines are usually implemented with the `sbrk(2)` system call

This system call expands (or contracts) the heap of the process

- However, most versions of `malloc` and `free` never decrease their memory size
- The space that we free is available for a later allocation
- The freed space is usually kept in the `malloc` pool, not returned to the kernel

# The alloca Function

---

A special memory allocation function – alloca

```
#include <alloca.h>
void *alloca(size_t size);
```

alloca() allocate memories in [stack frames](#) of the current function call

So you don't have to free() the memory – it is released automatically after the execution of the current function returns

May be not supported by your system, but modern UNIXes supports the function (Linux, FreeBSD, Mac OS X, Solaris)

Pros: might be faster (than malloc), no need to free, easier to work with setjmp/longjmp

Cons: Portability

# setjmp and longjmp Function

---

The reserved keyword "goto" can be used only in the same function

We cannot goto a label that is in another function

Instead, we must use the setjmp and longjmp functions to perform this type of branching

# Typical Program Skeleton for Command Processing

```
#include "apue.h"
```

```
#define TOK_ADD 5
```

```
void do_line(char *);  
void cmd_add(void);  
int get_token(void);
```

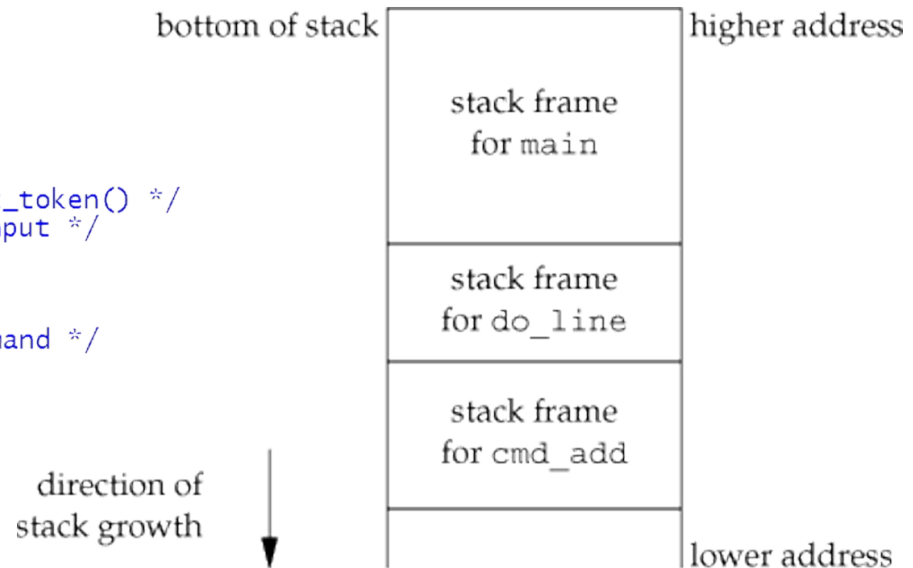
```
int main(void) {  
    char line[MAXLINE];  
    while (fgets(line, MAXLINE, stdin) != NULL)  
        do_line(line);  
    exit(0);  
}
```

```
char *tok_ptr; /* global pointer for get_token() */  
void do_line(char *ptr) { /* process one line of input */  
    int cmd;  
    tok_ptr = ptr;  
    while ((cmd = get_token()) > 0) {  
        switch (cmd) { /* one case for each command */  
            case TOK_ADD: cmd_add(); break;  
        }  
    }  
}
```

```
void cmd_add(void) {  
    int token;  
    token = get_token(); /* rest of processing for this command */  
}
```

```
int get_token(void) {  
    /* fetch next token from line pointed to by tok_ptr */  
}
```

What if we encounter an error in `cmd_add` and would like to jump back to the main function for processing the next line?



# The Solution for Jumping Across Functions

---

## Set the jump back position

- `int setjmp(jmp_buf env);`
- `env` is usually a global variable – has to be accessed from both the `setjmp` side and the `longjmp` side
- Returns: 0 if called directly, or nonzero if returning from a call to `longjmp`

## Jump back

- `void longjmp(jmp_buf env, int val);`
  - The 'val' will be returned from `setjmp`
  - If val is 0, it will be replaced by 1

# Using setjmp and longjmp

```
#include "apue.h"
#include <setjmp.h>

#define TOK_ADD    5


jmp_buf jmpbuffer;

int main(void) {
    char    line[MAXLINE];
    if (setjmp(jmpbuffer) != 0)
        printf("error");
    while (fgets(line, MAXLINE, stdin) != NULL)
        do_line(line);
    exit(0);
}

. . .

void cmd_add(void) {
    int      token;
    token = get_token();
    if (token < 0)          /* an error has occurred */
        longjmp(jmpbuffer, 1);
    /* rest of processing for this command */
}
```

Stack after jumped back



# Restoration of Variables (1/4)

---

## Type of variables

- Automatic, e.g., `[auto] int autoVal;`, the default
- Register, e.g., `register int regVal;`, store in register if possible
- Volatile, e.g., `volatile int volVal;`, store in memory

## What are the values of variables after jumped back?

- It depends
- Most implementations do not try to roll back these automatic variables and register variables
- The standards say only that their values are indeterminate
- If you have an automatic variable that you do not want to be rolled back, define it with the volatile attribute
- Variables that are declared global or static are left alone when `longjmp` is executed
- **In short: variables in register – restored; variables in memory – kept**

# Restoration of Variables (2/4)

```
#include "apue.h"
#include <setjmp.h>
static void f1(int, int, int, int);
static void f2(void);
static jmp_buf jmpbuffer;
static int globval;

int main(void) {
    int autoval;
    register int regival;
    volatile int volaval;
    static int statval;
    globval = 1; autoval = 2; regival = 3; volaval = 4; statval = 5;
    if (setjmp(jmpbuffer) != 0) {
        printf("after longjmp:\n");
        printf("globval = %d, autoval = %d, regival = %d,"
               " volaval = %d, statval = %d\n",
               globval, autoval, regival, volaval, statval);
        exit(0);
    }
    // Change variables after setjmp, but before longjmp.
    globval = 95; autoval = 96; regival = 97; volaval = 98; statval = 99;
    f1(autoval, regival, volaval, statval); /* never returns */
    exit(0);
}

static void f1(int i, int j, int k, int l) {
    printf("in f1():\n");
    printf("globval = %d, autoval = %d, regival = %d,"
           " volaval = %d, statval = %d\n", globval, i, j, k, l);
    f2();
}

static void f2(void) { longjmp(jmpbuffer, 1); }
```



# Restoration of Variables (3/4)

---

## Rules for variable restoration

- Variables stored in memory will have values as of the time of calling `longjmp`
- Variables in the CPU and floating-point registers are restored to their values when `setjmp` was called

Hence,

- auto variables may be indeterminate, it depends on compiler implementations
- register variables are restored to the value of “before calling **setjmp**”
- volatile variable are restored to the value of “before calling **longjmp**”

# Restoration of Variables (4/4)

---

Set 1,2,3,4,5 → setjmp → Set 95,96,97,98,99 → longjmp → ?

- No optimization: gcc places everything in memory
- Full optimization: auto/register variables are placed in registers

```
$ gcc fig7.13-testjmp.c -I../include -o t1      compile without any optimization
$ gcc fig7.13-testjmp.c -I../include -o t2 -O    compile with full optimization
$ ./t1
in f1():
globval = 95, autoval = 96, regival = 97, volaval = 98, statval = 99
after longjmp:
globval = 95, autoval = 96, regival = 97, volaval = 98, statval = 99
$ ./t2
in f1():
globval = 95, autoval = 96, regival = 97, volaval = 98, statval = 99
after longjmp:
globval = 95, autoval = 2, regival = 3, volaval = 98, statval = 99
```

# Process Resource Limits

---

Every process has a set of resource limits

Resource limits are usually initialized by a parent process and inherited by its child processes

The `getrlimit` and `setrlimit` functions

```
#include <sys/time.h>
#include <sys/resource.h>
int getrlimit(int resource, struct rlimit *rlim);
int setrlimit(int resource, const struct rlimit *rlim);
```

The `rlimit` structure

```
struct rlimit {
    rlim_t rlim_cur; /* Soft limit */
    rlim_t rlim_max; /* Hard limit (ceiling for rlim_cur) */
};
```

# Partial List of Process Resources

Limit	XSI	FreeBSD 8.0	Linux 3.2.0	Mac OS X 10.6.8	Solaris 10
RLIMIT_AS	•	•	•		•
RLIMIT_CORE	•	•	•	•	•
RLIMIT_CPU	•	•	•	•	•
RLIMIT_DATA	•	•	•	•	•
RLIMIT_FSIZE	•	•	•	•	•
RLIMIT_MEMLOCK		•	•	•	
RLIMIT_NOFILE	•	•	•	•	•
RLIMIT_NPROC		•	•	•	
RLIMIT_RSS		•	•	•	
RLIMIT_SBSIZE		•			
RLIMIT_STACK	•	•	•	•	•
RLIMIT_VMEM					•

# Example to Dump Resource Limits

---

See code fig7.16-getrlimit.c

```
$ ./fig7.16-getrlimit
RLIMIT_AS      (infinite) (infinite)
RLIMIT_CORE    1024000000 (infinite)
RLIMIT_CPU     (infinite) (infinite)
RLIMIT_DATA    (infinite) (infinite)
RLIMIT_FSIZE   (infinite) (infinite)
RLIMIT_LOCKS   (infinite) (infinite)
RLIMIT_MEMLOCK 65536      65536
RLIMIT_NOFILE  1024      4096
RLIMIT_NPROC   96120     96120
RLIMIT_RSS     (infinite) (infinite)
RLIMIT_STACK   8388608   (infinite)
```

# Example to Dump Resource Limits

---

Limits	Description
RLIMIT_CORE	The maximum size in bytes of a core file. A limit of 0 prevents the creation of a core file.
RLIMIT_MEMLOCK	The maximum amount of memory in bytes that a process can lock into memory using mlock(2).
RLIMIT_NOFILE	The maximum number of open files per process.
RLIMIT_NPROC	The maximum number of child processes per real user ID.
RLIMIT_STACK	The maximum size in bytes of the stack.

# Q & A

---