

# Threads

---

Advanced Programming in the UNIX Environment

Chun-Ying Huang <chuang@cs.nctu.edu.tw>

# Outline

---

Overview and introduction

Thread creation

Thread termination

Thread synchronization

# Introduction

---

We have introduced the relationships between processes

- There is only a limited amount shares can occur between related processes

Here we are going to introduce threads

- It is able to perform multiple tasks within a single process
- All threads within a single process have access to the same process components, e.g., file descriptors and memory

If a single resource is shared among multiple threads

- We need synchronization mechanisms to prevent multiple threads from viewing inconsistencies in their shared resources

# Thread Concepts

---

A typical UNIX process can be thought of as having a single thread of control

- Each process is doing only one thing at a time

With multiple threads of control

- We can design our programs to do more than one thing at a time within a single process
- Each thread handles a task

Benefits of using multiple threads

- Simplify code that deals with asynchronous events
- Shares the same memory spaces and file descriptors
- Problems can be partitioned to improve overall program throughput
- Interactive programs can realize improved response time

# Thread Concepts (Cont'd)

---

Multithread programming can be realized also on a single processor

- The performance still gets improved as there is always I/O operations that block the execution of a process
- However, if you have a multiprocessor system, threaded program may run faster

A thread consists of the information necessary to represent an execution context within a process

- Thread ID
- Register values, stack content, a signal mask, an errno variable
- Scheduling priority and policy, and
- Thread specific data

# The UNIX Thread Standard

---

## Defined by POSIX

- Portable Operating System Interface for UNIX
- POSIX.1-2001
- Also known as “pthreads” or “POSIX threads”

## Build programs with thread supports

- Your program has to include `<pthread.h>`
- To compile a C/C++ program with thread support, you have to add “-pthread” or “-lpthread” argument when compiling with gcc

# Linux Implementation of POSIX Threads

---

Threads are implemented via the clone system call

- Basically, they are processes sharing more information

Two flavors: LinuxThreads and NPTL

LinuxThreads: The original thread implementation

NPTL: Native POSIX Thread Library

- Better conformance to POSIX.1
  - For example, POSIX.1 requires threads of a process obtaining the same PID value when calling `getpid()`, but LinuxThreads does not follow it
- Better performance
- Require supports from the C library and the kernel

Both are 1:1 thread model

That is, each thread maps to a kernel scheduling entity

# Thread Identification

---

Every thread has a thread ID

- It may be not unique in the system
- The thread ID has significance only within the context of the process to which it belongs

The `pthread_t` data type

- Similar to `pid_t`, `pthread_t` is used to identify a thread
- It can be a structure (not forced to be an integer)

Test the equivalence of thread IDs

- `int pthread_equal(pthread_t tid1, pthread_t tid2);`
- Returns: nonzero if equal, zero otherwise

Get the current thread ID

- `pthread_t pthread_self(void);`



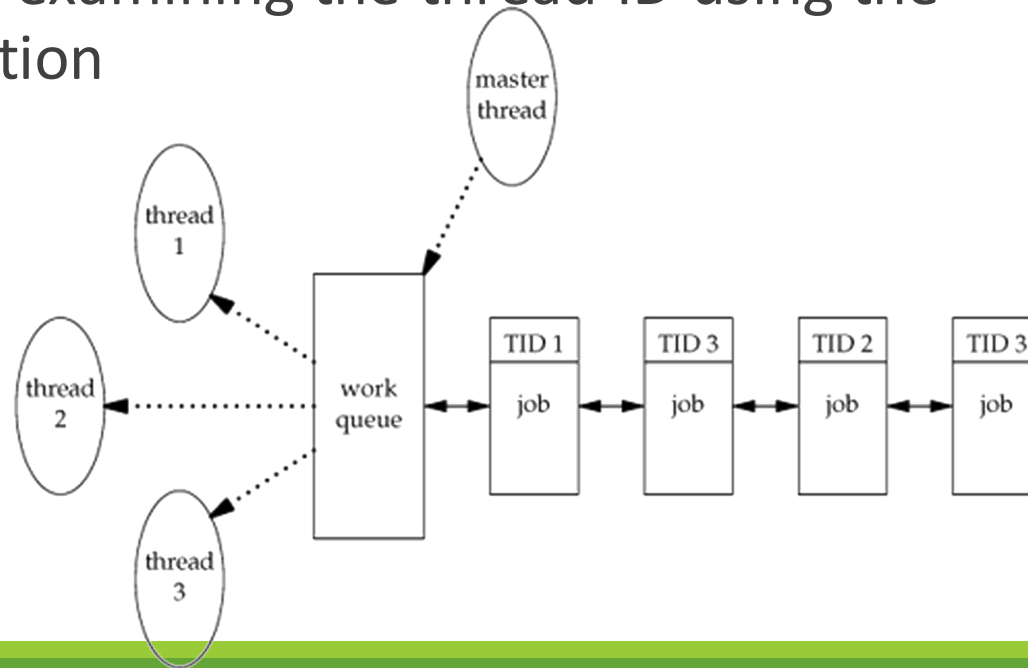
# Thread ID: A Job Queue Example

---

A master thread assign jobs to worker threads by their IDs

A worker thread only removes the job tagged with its own thread ID

This can be done by examining the thread ID using the `pthread_equal` function



# Thread Creation

---

With pthreads, when a program runs, it also starts out as a single process with a single thread of control

Create additional threads

- `int pthread_create(pthread_t *restrict thread, const pthread_attr_t *restrict attr, void *(*start_routine)(void*), void *restrict arg);`
  - “thread” should be the address of a previous declared `pthread_t` variable
  - “attr” is used to customize thread attributes
  - The newly created thread starts running the “start\_routine” function
  - The “arg” is then passed to the “start\_routine” function
- Returns: 0 if OK, error number on failure

# Thread Creation (Cont'd)

---

When a thread is created ...

- There is no guarantee which thread runs first
  - The newly created thread or the calling thread?
- The newly created thread
  - Has access to the process address space, and
  - Inherit floating-point environment and signal mask from the calling process

pthread functions usually return an error code when they fail

- They do not use the errno variable
- It is not recommended to use global variables for status checks
- However, the per thread copy of errno is still provided for compatibility

# Thread Creation, an Example

---

```
pthread_t ntid;
void printids(const char *s) {
    pid_t pid = getpid();
    pthread_t tid = pthread_self();
    tid = printf("%s pid %u tid %u (0x%x)\n", s, (unsigned int)pid,
        (unsigned int) tid, (unsigned int) tid);
}
void * thr_fn(void *arg) {
    printids("new thread: "); return((void *)0);
}
int main(void) {
    int err = pthread_create(&ntid, NULL, thr_fn, NULL);
    if (err != 0)
        err_quit("can't create thread: %s\n", strerror(err));
    printids("main thread:");
    sleep(1);
    exit(0);
}
```

# Thread Creation, an Example (Cont'd)

---

The result can be different on various platforms

- The `pthread_t` may be not an integer
- The `getpid()` function may return different values for the two threads (although it is expected to return the same value)

```
$ ./fig11.2-threadid  
new thread:  pid 3207 tid 3084950416 (0xb7e09b90)  
main thread: pid 3207 tid 3084953264 (0xb7e0a6b0)
```

# Thread Termination

---

## Terminate the entire process

- If any thread within a process calls `exit`, `_Exit`, or `_exit`
- If received a signal with the default action of terminating the process

## Terminate a single thread

- Return from the start routine.
  - The return value is the thread's exit code
- Cancelled by another thread in the same process
- The thread calls `pthread_exit`

# Thread Termination Status

---

The exit status of a process can be retrieved using wait functions – wait, waitpid, ..., etc

The exit status of a thread can also be retrieved

The pthread\_join function

- `int pthread_join(pthread_t thread, void **value_ptr);`
- Returns: 0 if OK, error number on failure
- This function **suspends** the calling thread
  - Unless the target thread has already terminated
- The retrieved exit status is stored in value\_ptr, if it is not NULL
- The target thread is then placed in a “detached” state
  - The storage for that thread is released

# Thread Termination Status (Cont'd)

---

The storage of a thread can be released immediately right on its termination

The `pthread_detach` function

- Set the state of a thread to be “detached”
- `int pthread_detach(pthread_t thread);`
- Returns: 0 if OK, error number on failure

A detached thread can not be joined

- A call to `pthread_join` for a detached thread will return `EINVL`



# Thread Termination, an Example

---

```
void * tfn1(void *a) { printf("thread 1 returning\n"); return((void *)1); }
void * tfn2(void *a) { printf("thread 2 exiting\n"); pthread_exit((void *)2); }
int main(void) {
    int err;
    pthread_t tid1, tid2;
    void *tret;
    err = pthread_create(&tid1, NULL, tfn1, NULL);
    if (err != 0)err_quit("can't create thread 1: %s\n", strerror(err));
    err = pthread_create(&tid2, NULL, tfn2, NULL);
    if (err != 0)err_quit("can't create thread 2: %s\n", strerror(err));
    err = pthread_join(tid1, &tret);
    if (err != 0)err_quit("can't join with thread 1: %s\n", strerror(err));
    printf("thread 1 exit code %d\n", (int)tret);
    err = pthread_join(tid2, &tret);
    if (err != 0)err_quit("can't join with thread 2: %s\n", strerror(err));
    printf("thread 2 exit code %d\n", (int)tret);
    exit(0);
}
```

```
$ ./fig11.3-exitstatus
thread 1 returning
thread 2 exiting
thread 1 exit code 1
thread 2 exit code 2
```

# void \* and pthread Functions

---

In `pthread_create` and `pthread_exit` function, we pass arguments using the "void \*" type

- The typeless pointer

The pointer can be used to pass more than a single value

- Values can be store in a data structure
- The pointer of the data structure is then passed to `pthread_create` or `pthread_exit`

However, the data structure should not be placed on the stack

- When a thread is terminated, the memory of its stack is released
- It might be reused by other threads

# Cancelling a Thread

---

The `pthread_cancel` function

- `int pthread_cancel(pthread_t tid);`
- Returns: 0 if OK, error number on failure

It just like the thread *tid* calls  
`pthread_exit(PTHREAD_CANCELED)`

The thread *tid* can select to ignore or control how it is canceled

`pthread_cancel` does not wait for the thread to terminate

- It simply makes the request

# Cleanup Functions

---

Recall the atexit function

- Register functions that execute when a process terminates

Similar works can be done for threads

- `void pthread_cleanup_push(void (*rtn)(void *), void *arg);`
- `void pthread_cleanup_pop(int execute);`

The registered routines is executed when ...

- Making a call to `pthread_exit`
- Responding to a cancellation request
- Making a call to `pthread_cleanup_pop` with a nonzero execute argument
  - If the argument is zero, it just remove the routine on stack top

# Comparison of Process and Thread Primitives

---

Process Primitive	Thread Primitive	Description
fork	pthread_create	create a new flow of control
exit	pthread_exit	exit from an existing flow of control
waitpid	pthread_join	get exit status from flow of control
atexit	pthread_cleanup_push	register function to be called at exit from flow of control
getpid	pthread_self	get ID for flow of control
abort	pthread_cancel	request abnormal termination of flow of control

# Thread Synchronization

---

Threads of a process share the same memory

Each thread must see a consistent view of data

The data is always consistent if ...

- Each thread uses variables that other threads do not read or modify
- Variables are read-only

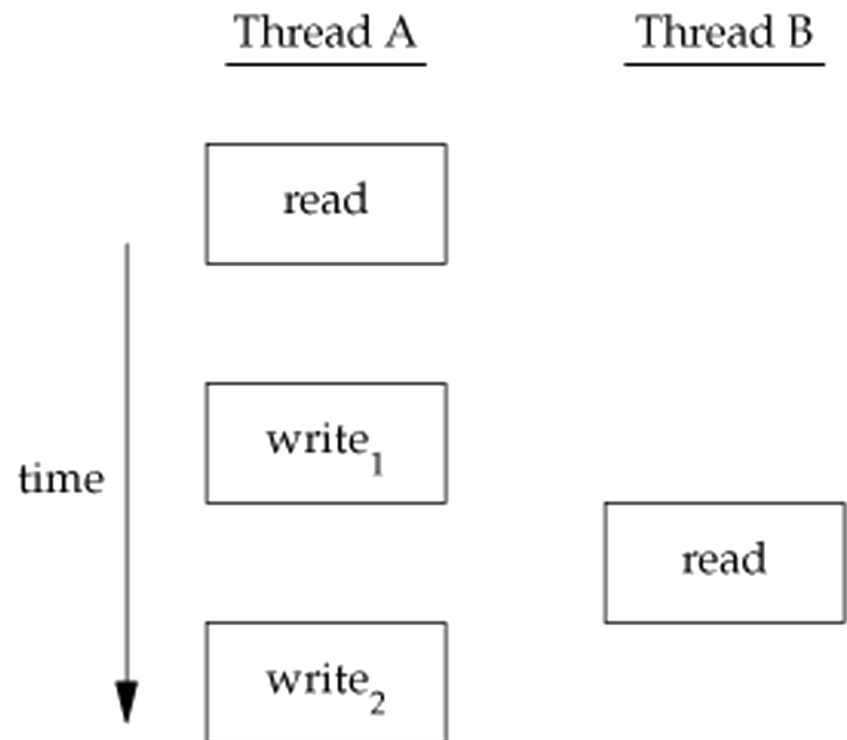
However, if a thread modifies a shared data

- We need to synchronize the threads to ensure that they do not use an invalid value

# Unsafe Access of Shared Variables, Example #1

Two threads, one for updating  
and one for reading

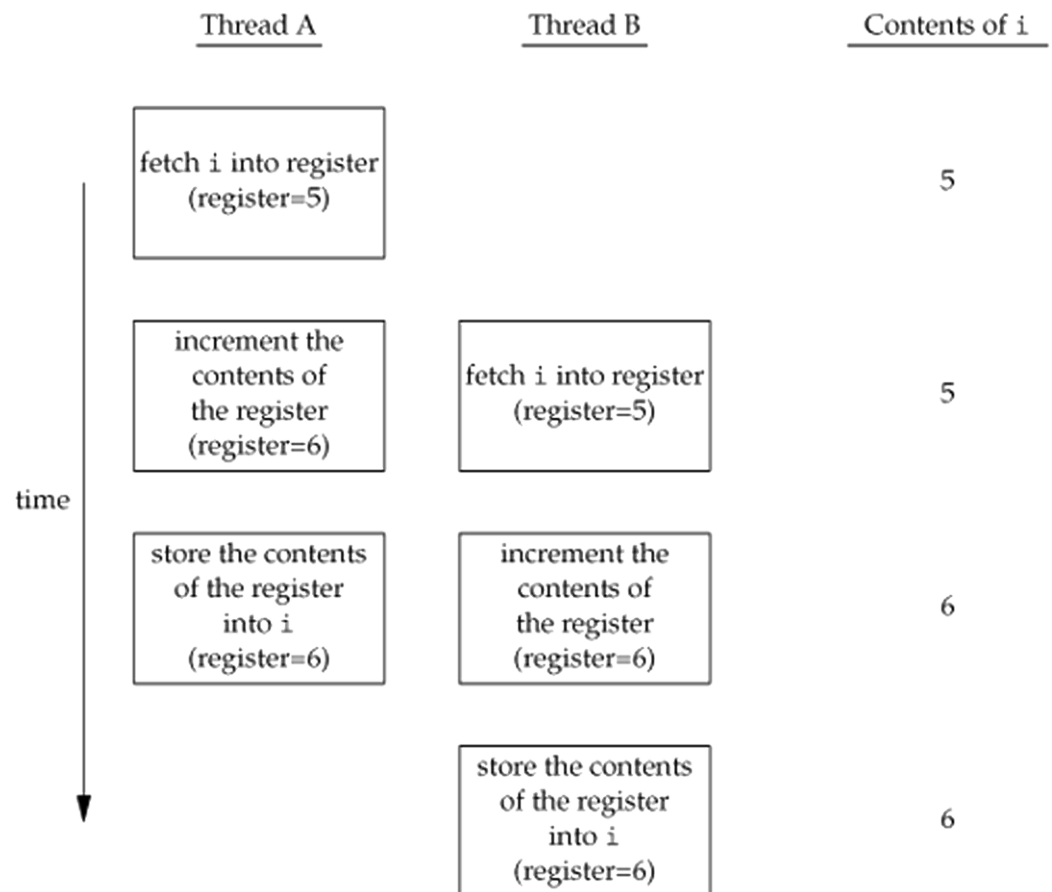
- Suppose a write operation needs two cycles
- The read operation occurs during the write operations



# Unsafe Access of Shared Variables, Example #2

Two threads, both  
increasing a variable

- Read the memory location into a register
- Increment the value in the register
- Write the new value back to the memory location





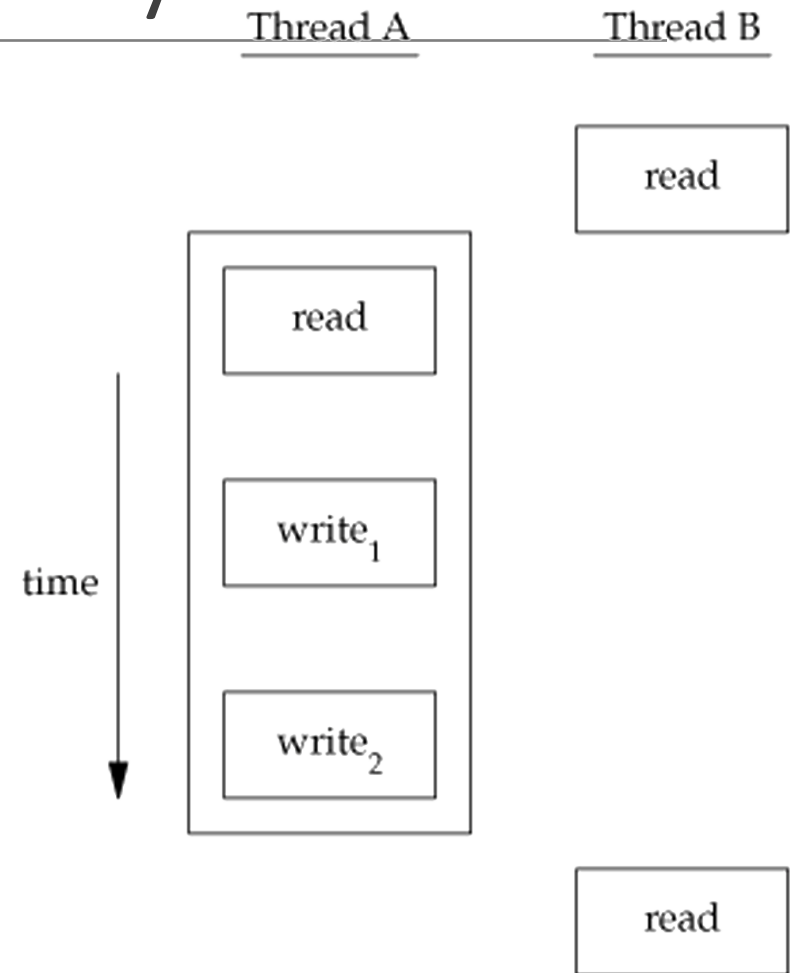
# Synchronized Memory Access

To solve the previous problem, we have to use a **lock** that allows only one thread to access the variable at a time

If thread B wants to read the variable, it acquires a lock

If thread A updates the variable, it acquires the same lock

- Thread B will be unable to read the variable until thread A releases the lock



# Mutexes

---

## Mutual exclusives

A mutex is basically a lock

- We set (lock) it before accessing a shared resource
- It is released (unlocked) when we're done

When a mutex is set ...

- Any other thread that tries to set it will be blocked until the lock holder releases it
- If more than one thread is blocked when a mutex is unlocked
  - All threads blocked on the lock will be made runnable
  - The first one to run will be able to set the lock
  - The others will see that the mutex is still locked and go back to wait
- Only one thread will proceed at a time

# pthread Mutexes

---

Data type: `pthread_mutex_t`

Initialize and destroy

- `int pthread_mutex_init(pthread_mutex_t* restrict mutex, const pthread_mutexattr_t* restrict attr);`
- `int pthread_mutex_destroy(pthread_mutex_t* mutex);`
- Returns: 0 if OK, error number on failure

Alternatively

- `pthread_mutex_t mutex = PTHREAD_MUTEX_INITIALIZER;`

# pthread Mutexes (Cont'd)

---

## Lock and unlock

- `int pthread_mutex_lock(pthread_mutex_t *mutex);`
- `int pthread_mutex_trylock(pthread_mutex_t *mutex);`
- `int pthread_mutex_unlock(pthread_mutex_t *mutex);`
- Returns: 0 if OK, error number on failure

# A Mutex Example – Protect a Data Structure

---

```
struct foo {
    int      f_count;
    pthread_mutex_t f_lock;
    /* ... more stuff here ... */
};

struct foo * foo_alloc(void) { /* allocate the object */
    struct foo *fp;
    if ((fp = malloc(sizeof(struct foo))) != NULL) {
        fp->f_count = 1;
        if (pthread_mutex_init(&fp->f_lock, NULL) != 0) {
            free(fp);
            return(NULL);
        }
        /* ... continue initialization ... */
    }
    return(fp);
}
```

# A Mutex Example – Protect a Data Structure (Cont'd)

---

```
void foo_hold(struct foo *fp) { /* add a reference to the object */
    pthread_mutex_lock(&fp->f_lock);
    fp->f_count++;
    pthread_mutex_unlock(&fp->f_lock);
}

void foo_rele(struct foo *fp) { /* release a reference to the object */
    pthread_mutex_lock(&fp->f_lock);
    if (--fp->f_count == 0) { /* last reference */
        pthread_mutex_unlock(&fp->f_lock);
        pthread_mutex_destroy(&fp->f_lock);
        free(fp);
    } else {
        pthread_mutex_unlock(&fp->f_lock);
    }
}
```

# Deadlock Avoidance

---

## How deadlock happens?

- Case #1: A thread lock the same mutex twice
- Case #2: Two threads (T1/T2) and two mutexes (MA/MB)
  - T1 locks MA and then locks MB
  - T2 locks MB and then locks MA
  - T1 and T2 may block each other

## Avoidance

- Case #1 is easier to avoid
- Case #2: Mutexes has to be locked in the same order
  - Every thread locks MA first and then locks MB
  - However, it is sometimes difficult to apply an ordered lock
- The `pthread_mutex_trylock` function
  - Make sure that we can lock all required mutexes at one time

# Reader-Writer Lock

---

Similar to mutexes, but allow higher degree of parallelism

With a mutex, it can be only

- Locked, or
- Unlocked

With a reader-write lock, it can be

- Locked in read mode
- Locked in write mode, or
- Unlocked

Reader-Write Lock

- Multiple reader locks can be acquired simultaneously
- Only one can lock in write mode
- If a reader/writer already locks, the coming writer/reader must wait until it unlocks



# pthread Reader-Writer Lock

---

Data type: `pthread_rwlock_t`

Initialize and destroy

- `int pthread_rwlock_init(pthread_rwlock_t *restrict rwlock, const pthread_rwlockattr_t *restrict attr);`
- `int pthread_rwlock_destroy(pthread_rwlock_t *rwlock);`
- Returns: 0 if OK, error number on failure

Lock and unlock

- `int pthread_rwlock_rdlock(pthread_rwlock_t *rwlock);`
- `int pthread_rwlock_wrlock(pthread_rwlock_t *rwlock);`
- `int pthread_rwlock_tryrdlock(pthread_rwlock_t *rwlock);`
- `int pthread_rwlock_trywrlock(pthread_rwlock_t *rwlock);`
- `int pthread_rwlock_unlock(pthread_rwlock_t *rwlock);`
- Returns: 0 if OK, error number on failure

# Condition Variable

---

Condition variable is another synchronization mechanism available to threads

It has to be used with mutexes

- The condition itself is protected by a mutex
- A thread must first lock the mutex to change the condition state

Condition variable allows a thread to wait in a race-free way for arbitrary conditions to occur

# pthread Condition Variables: Initialize and Destroy

---

Data type: `pthread_cond_t`

Initialize and destroy

- `int pthread_cond_init(pthread_cond_t *restrict cond, pthread_condattr_t *restrict attr);`
- `int pthread_cond_destroy(pthread_cond_t *cond);`
- Returns: 0 if OK, error number on failure

Alternatively

- `pthread_cond_t cond = PTHREAD_COND_INITIALIZER;`

# pthread Condition Variables: Wait for the Condition

---

## Synopsis

- `int pthread_cond_wait(pthread_cond_t *restrict cond, pthread_mutex_t *restrict mutex);`
- `int pthread_cond_timedwait(pthread_cond_t *restrict cond, pthread_mutex_t *restrict mutex, const struct timespec *restrict abstime);`

The condition wait function unlocks the mutex first

It then waits for the condition to occur

- The running state of the current thread is set to sleeping

# pthread Condition Variables: Timed Wait

---

The timespec data structure

It is the absolute time that the wait gives up

```
struct timespec {
    __time_t tv_sec;           /* Seconds. */
    long int tv_nsec;         /* Nanoseconds. */
};
```

An example of setting the absolute expire time

```
void maketimeout(struct timespec *tsp, long minutes) {
    struct timeval now; /* get the current time */
    gettimeofday(&now);
    tsp->tv_sec = now.tv_sec;
    tsp->tv_nsec = now.tv_usec * 1000; /* usec to nsec */
    /* add the offset to get timeout value */
    tsp->tv_sec += minutes * 60;
}
```

# pthread Condition Variables: Notifications

---

Notify threads that a condition has been satisfied

- `int pthread_cond_broadcast(pthread_cond_t *cond);`
- `int pthread_cond_signal(pthread_cond_t *cond);`
- Returns: 0 if OK, error number on failure

`pthread_cond_broadcast`

- Wake up all waiting threads

`pthread_cond_signal`

- Wake up one waiting threads
- POSIX.1 allows the implementation wakes up more than one threads
- Waked up threads have to contend for the mutex

# pthread Condition Variables: An Example

---

```
struct msg { struct msg *m_next; /* ... more stuff here ... */ };
struct msg *workq;
pthread_cond_t qready = PTHREAD_COND_INITIALIZER;
pthread_mutex_t qlock = PTHREAD_MUTEX_INITIALIZER;
void process_msg(void) {
    struct msg *mp;
    for (;;) {
        pthread_mutex_lock(&qlock);
        while (workq == NULL) { pthread_cond_wait(&qready, &qlock); }
        mp = workq;
        workq = mp->m_next;
        pthread_mutex_unlock(&qlock);
        /* now process the message mp */
    }
}

void enqueue_msg(struct msg *mp) {
    pthread_mutex_lock(&qlock);
    mp->m_next = workq;
    workq = mp;
    pthread_mutex_unlock(&qlock);
    pthread_cond_signal(&qready);
}
```

# Example: An Implementation of a Worker Queue – Jobs

---

## Job header

```
class Job {
private:
    pthread_t tid;
    int ch;
public:
    Job(int ch = 0, pthread_t tid = 0);
    pthread_t getId();
    void setId(pthread_t tid);
    int getChar();
    void setChar(int ch);
};
```

## Job Implementation

```
Job::Job(int ch, pthread_t tid) {
    this->ch = ch;
    this->tid = tid;
}
pthread_t Job::getId() {
    return tid;
}
void Job::setId(pthread_t tid) {
    this->tid = tid;
}
int Job::getChar() {
    return ch;
}
void Job::setChar(int ch) {
    this->ch = ch;
}
```



# Example: An Implementation of a Worker Queue – Global Definition

---

```
#define ASSIGNID      /* Assign thread id to jobs */
#define ORDERED      /* Ensure that jobs are processed in the order */
#define N_WORKERS      3

pthread_mutex_t mutex = PTHREAD_MUTEX_INITIALIZER;
pthread_cond_t cond = PTHREAD_COND_INITIALIZER;
std::list<Job> jobqueue;

int do_the_job(long id, int ch) {
    if(ch == -1) /* terminate the worker */
        return -1;
    printf("worker-%ld: %c\n", id, ch);
    return 0;
}
```

# Example: An Implementation of a Worker Queue – main Func (1/4)

---

```
int main(int argc, char *argv[]) {
    pthread_t workers[N_WORKERS];
    // check args
    if(argc < 2) {
        fprintf(stderr, "usage: %s input-string\n", argv[0]);
        return -1;
    }
    // create workers
    for(int i = 0; i < N_WORKERS; i++) {
        if(pthread_create(&workers[i], NULL,
                        worker_main, (void *) (long) i) != 0) {
            fprintf(stderr, "create worker[%d] failed\n", i);
            exit(-1);
        }
    }
}
```

# Example: An Implementation of a Worker Queue – main Func (2/4)

---

```
// create jobs
for(char *ptr = argv[1]; *ptr; ptr++) {
#ifdef ASSIGNID
    Job j(*ptr, workers[(ptr - argv[1]) % N_WORKERS]);
#else
    Job j(*ptr);
#endif
    pthread_mutex_lock(&mutex);
    jobqueue.push_back(j);
    pthread_mutex_unlock(&mutex);
    pthread_cond_signal(&cond);
}
```

# Example: An Implementation of a Worker Queue – main Func (3/4)

---

```
// terminate workers
for(int i = 0; i < N_WORKERS; i++) {
#ifdef ASSIGNID
    Job j(-1, workers[i]);
#else
    Job j(-1);
#endif
    pthread_mutex_lock(&mutex);
    jobqueue.push_back(j);
    pthread_mutex_unlock(&mutex);
    pthread_cond_signal(&cond);
}
```

# Example: An Implementation of a Worker Queue – main Func (4/4)

---

```
// process all jobs
size_t jobs;
do {
    pthread_mutex_lock(&mutex);
    jobs = jobqueue.size();
    pthread_mutex_unlock(&mutex);
    pthread_cond_signal(&cond);
} while(jobs > 0);
// wait for all workers
for(int i = 0; i < N_WORKERS; i++) {
    void *ret;
    pthread_join(workers[i], &ret);
}
//
return 0;
} /* end of main() */
```

# Example: An Implementation of a Worker Queue – The Worker

```
void* worker_main(void *arg) {
    long id = (long) arg;
    printf("# worker-%ld created\n", id);
    while(1) {
        Job j;
        pthread_mutex_lock(&mutex);
        pthread_cond_wait(&cond, &mutex);
        // has at least one job
        j = jobqueue.front();
        if(j.getId() == 0
           || pthread_equal(pthread_self(), j.getId())) {
            jobqueue.pop_front();
#ifdef ORDERED
            // will follow the order
            if(do_the_job(id, j.getChar()) < 0) {
                pthread_mutex_unlock(&mutex);
                break;
            }
#endif
        } else {
            pthread_mutex_unlock(&mutex);
            continue;
        }
        pthread_mutex_unlock(&mutex);
#ifdef ORDERED
        // may be out of order
        if(do_the_job(id, j.getChar()) < 0)
            break;
#endif
    }
    printf("# worker-%ld terminated\n",
           id);
    return NULL;
}
```

# Barriers

---

Barriers are used to coordinate multiple threads working in parallel

Allow each thread to wait until all cooperating threads have reached the same point

Create and destroy a barrier

```
int pthread_barrier_init(pthread_barrier_t *restrict barrier,  
                        const pthread_barrierattr_t *restrict attr,  
                        unsigned count);  
int pthread_barrier_destroy(pthread_barrier_t *barrier);
```

Wait for a barrier

```
int pthread_barrier_wait(pthread_barrier_t *barrier);
```

# Barrier Example (1/3)

---

```
#define HAS_BARRIER
#define N 5

static pthread_barrier_t barrier;

void *worker(void *arg) {
    long i, id = (long) arg;
    for(i = 0; i < id+1; i++) {
        printf("%ld", id+1);
    }
    printf("[%ld/done]\n", id+1);
#ifdef HAS_BARRIER
    pthread_barrier_wait(&barrier);
#endif
    return NULL;
}
```



# Barrier Example (2/3)

---

```
int main() {
    long i;
    pthread_t tid;
#ifdef HAS_BARRIER
    pthread_barrier_init(&barrier, NULL, N+1);
#endif
    for(i = 0; i < N; i++) {
        if(pthread_create(&tid, NULL, worker, (void*) i) != 0) {
            fprintf(stderr, "pthread_create failed.\n");
            return -1;
        }
    }
#ifdef HAS_BARRIER
    pthread_barrier_wait(&barrier);
    pthread_barrier_destroy(&barrier);
#endif
    printf("all done.\n");
    return 0;
}
```

# Barrier Example (3/3)

---

## Without HAS\_BARRIER

```
$ ./barrier
3314444[4/done]
all done.
```

```
$ ./barrier
all done.
```

```
$ ./barrier
333[3/done]
1[1/done]
22[2/done]
4444[4/done]
all done.
```

## With HAS\_BARRIER

```
$ ./barrier
22[2/done]
13[1/done]
4444[4/done]
55555[5/done]
33[3/done]
all done.
```

```
$ ./barrier
43233[3/done]
2[2/done]
555551[1/done]
[5/done]
444[4/done]
all done.
```

# Q & A

---