

Classical Inter-Process Communication

Advanced Programming in the UNIX Environment

Chun-Ying Huang <chuang@cs.nctu.edu.tw>

Introduction

We have described the process control primitives and seen how to invoke multiple processes

How does a process communicate with other processes?

The inter-process communication (IPC)

Common IPC Mechanisms

(Half-duplex) pipes

FIFOs

Message queues

Semaphores

Shared memory

Sockets

Pipes

The oldest form of UNIX System IPC

Historically, they have been **half duplex**

- Some modern system has full duplex pipe, but for program portability, it is not suggested to use full duplex pipe.

Pipes can be used only between processes that **have a common ancestor**

- Normally, a pipe is created by a process
- The process then calls fork
- The pipe is then used between the parent and the child

Creating a Pipe


Synopsis

```
int pipe(int filedes[2]);
```

- Returns: 0 if OK, -1 on error

Two descriptors are created

- `filedes[0]` is opened for reading, and
- `filedes[1]` is opened for writing




Sharing a Pipe

A pipe in a single process is useless

Normally, the process that calls pipe then calls fork


- This creates an IPC channel from the parent to the child or vice versa



Sharing a Pipe (Cont'd)

As the pipe is half duplex, the following actions may apply depending on the scenario

- If the pipe is used for a child to send data to its parent
 - The parent closes `fd[1]` and the child closes `fd[0]`
- If the pipe is used for a parent to send data to its child
 - The parent closes `fd[0]` and the child closes `fd[1]`, see the figure



An Example of Creating a Pipe

```
int main(void) {
    int n;
    int fd[2];
    pid_t pid;
    char line[MAXLINE];
    if (pipe(fd) < 0)
        err_sys("pipe error");
    if ((pid = fork()) < 0) {
        err_sys("fork error");
    } else if (pid > 0) {                                /* parent */
        close(fd[0]);
        write(fd[1], "hello world\n", 12);
    } else {                                            /* child */
        close(fd[1]);
        n = read(fd[0], line, MAXLINE);
        write(STDOUT_FILENO, line, n);
    }
    exit(0);
}
```


Process Synchronization: Using a Pipe

Recall: In Chapter 8

- Race Conditions between the Parent and the Child

```
int main(void) {
    pid_t  pid;
+    TELL_WAIT();
    if ((pid = fork()) < 0)      {
        err_sys("fork error");
    } else if (pid == 0) {
+        WAIT_PARENT();          /* parent goes first */
        charatime("output from child\n");
    } else {
        charatime("output from parent\n");
+        TELL_CHILD(pid);
    }
    exit(0);
}
```


Process Synchronization: Using a Pipe (Cont'd)

```
static int pfd1[2], pfd2[2];

void TELL_WAIT(void) {
    if (pipe(pfd1) < 0 || pipe(pfd2) < 0)
        err_sys("pipe error");
}

void WAIT_PARENT(void) {
    char c;
    if (read(pfd1[0], &c, 1) != 1)
        err_sys("read error");
    if (c != 'p')
        err_quit("WAIT_PARENT: incorrect data");
}

void TELL_CHILD(pid_t pid) {
    if (write(pfd1[1], "p", 1) != 1)
        err_sys("write error");
}
```



popen and pclose Functions

Execute a command and access its standard I/O

- Read from its standard output, or
- Write to its standard input
- As we are using half-duplex pipe, we cannot read/write at the same time

Synopsis

- `FILE *popen(const char *cmdstring, const char *type);`
- Returns: file pointer if OK, NULL on error
- `int pclose(FILE *fp);`
- termination status of *cmdstring*, or -1 on error

popen and pclose Functions


Operations

- create a pipe (pipe)
- fork a child (fork)
- close the unused ends of the pipe (close)
- configure the descriptor (dup2)
- execute a shell to run the command (exec), and
- wait for the command to terminate (wait)

popen with a type of “r”



popen with a type of “w”



Implementation of popen and pclose

See ch15/fig15.12-popen.c, or textbook figure 15.12

popen

- Make sure that type is “r” or “w”
- Create a buffer for popen children PIDs
- Create a pipe and fork a child process
- For the child:
 - If type is “r”, close fd[0], otherwise close fd[1]
 - `execl("/bin/sh", "sh", "-c", cmdstring, (char *)0);`
- For the parent
 - If type is “r”, close fd[1], otherwise close fd[0]
 - If type is “r”, `FILE *fp = fdopen(fd[0], type)`
 - Otherwise, `FILE *fp = fdopen(fd[1], type)`
 - Save child PID (indexed by pipe fd) and return fp

Implementation of popen and pclose (Cont'd)

pclose

- Get descriptor number by `fd = fileno(fp);`
- Retrieve the pid (indexed by pipe fd)
- Reset the corresponding pid on the children's pid buffer to zero
- `fclose(fp)`
- `waitpid(pid, &stat, 0)`
- `return(stat)`

popen Example: Filters


A filter that converts uppercases into lowercases

```
int main(void) {
    int c;
    while ((c = getchar()) != EOF) {
        if (isupper(c))
            c = tolower(c);
        if (putchar(c) == EOF)
            err_sys("output error");
        if (c == '\n')
            fflush(stdout);
    }
    exit(0);
}
```

popen Example: Filters (Cont'd)

A program that run the filter using popen, and show the filtered content

```
int main(void) {
    char    line[MAXLINE];
    FILE    *fpin;
    if ((fpin = popen("./myucl", "r")) == NULL)
        err_sys("popen error");
    for ( ; ; ) {
        fputs("prompt> ", stdout);
        fflush(stdout);
        if (fgets(line, MAXLINE, fpin) == NULL) /* read from pipe */
            break;
        if (fputs(line, stdout) == EOF)
            err_sys("fputs error to pipe");
    }
    if (pclose(fpin) == -1)
        err_sys("pclose error");
    putchar('\n');
    exit(0);
}
```




Coprocess

Definition of an UNIX system filter

- A process that reads from standard input and writes to standard output

Coprocess

- An UNIX system filter becomes a coprocess if the filter's input and output are both associated with the same program
- We need two pipe() calls to setup the communication channel between a program and its coprocess



Coprocess, an Example

A filter that read from STDIN, adds two numbers, and write to STDOUT

- Implemented using file I/O

```
int main(void) {
    int    n, int1, int2;
    char    line[MAXLINE];
    while ((n = read(STDIN_FILENO, line, MAXLINE)) > 0) {
        line[n] = 0                                /* null terminated */
        if (sscanf(line, "%d%d", &int1, &int2) == 2) {
            sprintf(line, "%d\n", int1 + int2);
            n = strlen(line);
            if (write(STDOUT_FILENO, line, n) != n)
                err_sys("write error");
        } else {
            if (write(STDOUT_FILENO, "invalid args\n", 13) != 13)
                err_sys("write error");
        }
    }
    exit(0);
}
```

The `sig_pipe` function
just print a message and
then `exit(1)`;

Coprocess, an Example (Cont'd)

```
int main(void) {
    int n, fd1[2], fd2[2];
    pid_t pid;
    char line[MAXLINE];
    if (signal(SIGPIPE, sig_pipe) == SIG_ERR)
        err_sys("signal error");
    if (pipe(fd1) < 0 || pipe(fd2) < 0)
        err_sys("pipe error");
    if ((pid = fork()) < 0) err_sys("fork error");
    else if (pid > 0) { /* parent */
        close(fd1[0]);
        close(fd2[1]);
        while (fgets(line, MAXLINE, stdin) != NULL) {
            n = strlen(line);
            if (write(fd1[1], line, n) != n)
                err_sys("write error to pipe");
            if ((n = read(fd2[0], line, MAXLINE)) < 0)
                err_sys("read error from pipe");
            if (n == 0) {
                err_msg("child closed pipe");
                break;
            }
            line[n] = 0; /* null terminate */
            if (fputs(line, stdout) == EOF)
                err_sys("fputs error");
        }
    }
```

```
        if (ferror(stdin))
            err_sys("fgets error on stdin");
        exit(0);
    } else { /* child */
        close(fd1[1]);
        close(fd2[0]);
        if (fd1[0] != STDIN_FILENO) {
            if (dup2(fd1[0], STDIN_FILENO) != STDIN_FILENO)
                err_sys("dup2 error to stdin");
            close(fd1[0]);
        }
        if (fd2[1] != STDOUT_FILENO) {
            if (dup2(fd2[1], STDOUT_FILENO) != STDOUT_FILENO)
                err_sys("dup2 error to stdout");
            close(fd2[1]);
        }
        if (execl("./add2", "add2", (char *)0) < 0)
            err_sys("execl error");
    }
    return 0;
}
```

Coprocess and Standard I/O

What happens if the coprocess is implemented using standard I/O?

- The filter no longer works!

It is because the I/O buffering mode

- When standard input/output are not terminal devices, they are fully buffered
- Solution: We need pseudo-terminals devices to emulate the line buffer or unbuffered channel (not discussed in this Chapter)

FIFOs

First in, first out

FIFOs are sometimes called named pipes

Pipes can be only used between processes of a common ancestor

With FIFOs, unrelated processes can exchange data

Creating a FIFO, synopsis

- `int mkfifo(const char *pathname, mode_t mode);`
- Returns: 0 if OK, -1 on error

Once we have used `mkfifo` to create a FIFO, we open it using `open`

Open an FIFO

When we open a FIFO, the non-blocking flag (`O_NONBLOCK`) affects what happens

In the normal case (`O_NONBLOCK` not specified)

- An open for read-only blocks until some other process opens the FIFO for writing
- Similarly, an open for write-only blocks until some other process opens the FIFO for reading

If `O_NONBLOCK` is specified

- An open for read-only returns immediately
- But an open for write-only returns -1 with `errno` set to `ENXIO` if no process has the FIFO open for reading

Share an FIFO

It is common to have multiple writers for a given FIFO

We have to worry about atomic writes if we don't want the writes from multiple processes to be interleaved

Applications of FIFOs

Data passing

- Pass data without creating intermediate temporary files


Client-server communication

- Used as rendezvous points in client-server applications

FIFO Applications – Data Passing

Scenario


- Process a filtered input stream twice



FIFO Applications – Data Passing (Cont'd)

Solutions with FIFO

- ```
$ mkfifo fifo1
$ prog3 < fifo1 &
$ prog1 < infile | tee fifo1 | prog2
```




# FIFO Applications – Client-Server Communication

---

## Scenario #1: One way communication

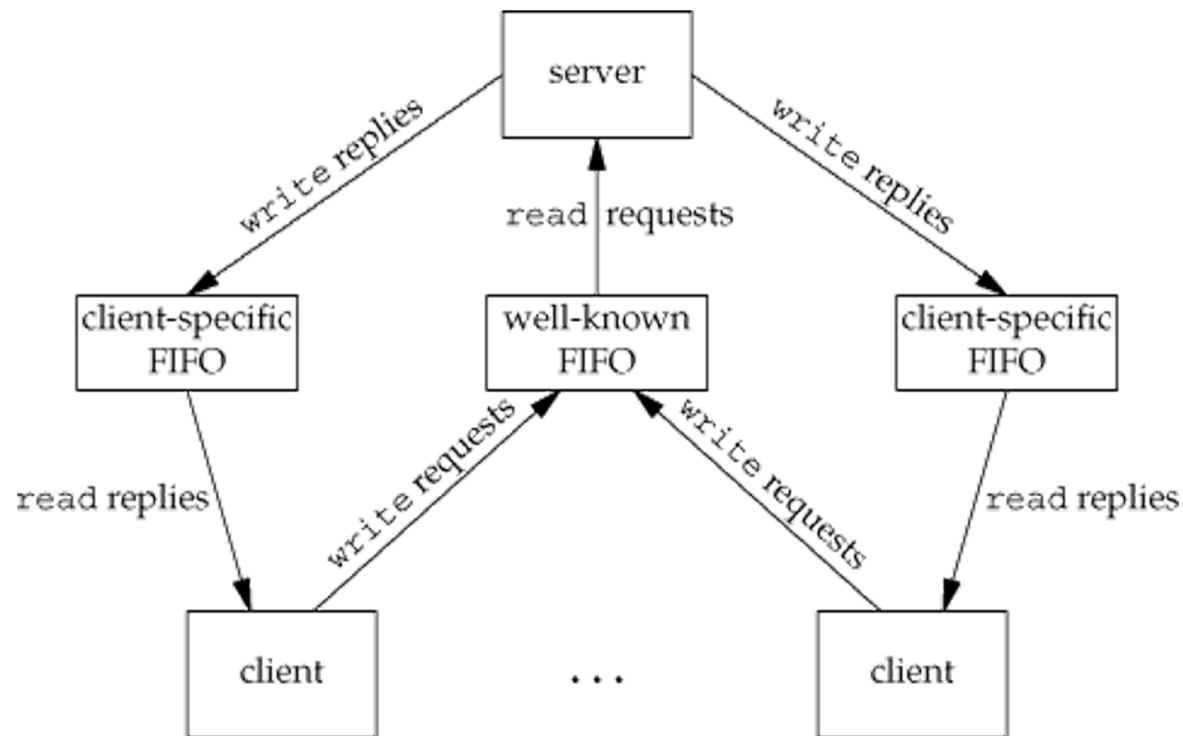
- Clients send requests to a server



# FIFO Applications – Client-Server Communication (Cont'd)

## Scenario #2: Two-way communications

- Client-server communication using FIFOs



# XSI (SysV) IPC

---

XSI – X/Open System Interface

Three types of XSI IPC

- Message queue
- Semaphore
- Shared memory

Common user commands

- `ipcs` – list IPC objects
- `ipcrm` – remove IPC objects

# XSI (SysV) IPC (Cont'd)

---

## IPC identifiers

- Each IPC structure in the kernel is referred to by a non-negative integer identifier
- We need to know the identifier to access the IPC object

However, the identifier is an internal name for an IPC object

- We need a naming scheme to refer the same IPC object – the IPC keys

## IPC keys

- Whenever an IPC structure is being created , a key must be specified
- Keys are of data type `key_t`
- Then, the identifier of the referred IPC object is returned

# Sharing of IPC Objects

---

A server can create an IPC object with a key of `IPC_PRIVATE`

- The identifier of the created IPC object can be passed by storing in a file, or
- Fork a child, which inherits the identifier directly

A server and a client can agree on a key by defining the key in a common header

A server and a client can agree on a pathname and a project ID

- The key can be generated by the `ftok` function
- `key_t ftok(const char *path, int id);`
- `path` must be an existing file, and
- `id` is a 8-bit non-zero number (you can not use more than 8 bits!)

# XSI IPC – Advantages and Disadvantages

---

## Advantages

- Reliable
- Supports flow control
- Record based
- Can be processed in other than first-in, first-out order

## Disadvantages

- IPC data may left in the system even if no one refers to it
- They are different from file system objects, i.e. no descriptors
- Therefore, we need a different set of system calls to manipulate them



# Message Queues

---

A message queue is a linked list of messages stored within the kernel

Each queue has a message queue identifier

Creating or opening a message queue

- `int msgget(key_t key, int flag);`
- Returns: 0 if OK, -1 on error
- Upon creating, the least significant 9 bits of *flag* define the permissions for the message queue
- *flag* can be OR'ed with `IPC_CREAT` and/or `IPC_EXCL`

# Message Queue – System Limitations

---

The limitations may vary on different platforms

- “ipcs -l” command on Linux
- “ipcs -Q” on BSD and Mac OS X

```
$ ipcs -l
```

```
...
```

```
----- Messages Limits -----
```

```
max queues system wide = 32768
```

```
max size of message (bytes) = 8192
```

```
default max size of queue (bytes) = 16384
```

# Controlling a Message Queue

---

The internal data structure associated with a message queue

```
struct msqid_ds {
 struct ipc_perm msg_perm; /* Ownership and permissions */
 time_t msg_stime; /* Time of last msgsnd(2) */
 time_t msg_rtime; /* Time of last msgrcv(2) */
 time_t msg_ctime; /* Time of last change */
 unsigned long __msg_cbytes; /* Current number of bytes in queue (non-standard) */
 msgqnum_t msg_qnum; /* Current number of messages in queue */
 msglen_t msg_qbytes; /* Maximum number of bytes allowed in queue */
 pid_t msg_lspid; /* PID of last msgsnd(2) */
 pid_t msg_lrpid; /* PID of last msgrcv(2) */
};
```

# Controlling a Message Queue (Cont'd)

---

## Synopsis

- `int msgctl(int msqid, int cmd, struct msqid_ds *buf);`
- Returns: 0 if OK, -1 on error

## The *cmd* can be

- `IPC_STAT`: Retrieve the internal `msqid_ds` data structure
- `IPC_SET`: Set the `msqid_ds`
  - `msg_perm.uid`, `msg_perm.gid`, `msg_perm.mode`, and `msg_qbytes`
  - Only superuser is able to increase `msg_qbytes`
- `IPC_RMID`: Remove the queue (immediately)

# Send a Message into Queue

---

## Synopsis

- `int msgsnd(int msqid, const void *ptr, size_t nbytes, int flag);`

The message, which is pointed to by *ptr*

- It must be started with a long integer (the type of the message)
- A *nbytes* message follows the long integer

```
struct msgbuf {
 long type; /* message type, must be > 0 */
 char mtext[1]; /* message data */
};
```

- The flag
  - `IPC_NOWAIT`: non-blocking access to the queue
  - If the queue is full and `IPC_NOWAIT` is specified
    - It returns an error with `errno` set to `EAGAIN`

# Receive a Message from Queue

---

## Synopsis

- `ssize_t msgrcv(int msqid, void *ptr, size_t nbytes, long type, int flag);`
- Returns: size of data portion of message if OK, -1 on error

## The message type

- If `type == 0`, *the first message* on the queue is returned
- If `type > 0`, *the first message* on the queue whose message type *equals* type is returned
- If `type < 0`, the first message on the queue whose message type is *the lowest value less than or equal to* the absolute value of type is returned

# Receive a Message from Queue (Cont'd)

---

## The flags

- IPC\_NOWAIT: non-blocking access to the queue
- MSG\_EXCEPT
  - If type > 0, *the first message* on the queue whose message type *has a non-equal* type is returned
- MSG\_NOERROR
  - If the received message has a longer size than *nbytes*, it is *truncated* and then returned

# Message Queue: Hello, World!

## Example

---

```
struct msgbuf {
 long mtype; /* message type, must be > 0 */
 char mtext[0]; /* message data */
};

int main() {
 int qid = -1, rlen, wlen;
 char buf[1024];
 pid_t pid;
 struct msgbuf *msg = (struct msgbuf*) buf;
 //
 if((qid = msgget(IPC_PRIVATE, IPC_CREAT|IPC_EXCL|0660)) < 0)
 err_sys("msgget");
 if((pid = fork()) < 0)
 err_sys("fork");
```



# Message Queue: Hello, World!

## Example (Cont'd)

---

```
if(pid == 0) { /* child */
 msg->mtype = 0;
 if((rlen = msgrcv(qid, msg, sizeof(buf)-sizeof(*msg), 0, 0)) < 0)
 err_sys("msgrcv");
 printf("[%ld] %s (%u bytes)\n", msg->mtype, msg->mtext, rlen);
} else { /* parent */
 msg->mtype = 1024;
 wlen = snprintf(msg->mtext, sizeof(buf)-sizeof(*msg),
 "%s", MESSAGE);
 if(msgsnd(qid, msg, wlen+1, 0) < 0)
 perror("msgsnd");
 else if(wait(&wlen) < 1)
 perror("wait");
 if(qid >= 0)
 if(msgctl(qid, IPC_RMID, NULL) < 0)
 err_sys("msgctl(RMID)");
}
return 0;
}
```

# Semaphore (1/3)

---

A semaphore is a shared counter

It is used to provide access to a shared data object for multiple processes

Procedures for a process to obtain a shared resource

- Test the semaphore that controls the resource
- If the value of the semaphore is positive, the process can use the resource
  - The process decrements the semaphore value by 1
- If the value of the semaphore is 0
  - The process goes to sleep until the semaphore value is greater than 0

# Semaphore (2/3)

---

## Features

- A semaphore is a set of one or more semaphore values
  - It is not simply a single non-negative value
- Independent of semaphore creation (semget) and initialization (semctl)
  - It may be a problem as we cannot *atomically* create a new semaphore set and initialize all the values in the set
- All XSI IPC objects are not released automatically
  - They remain in existence even when no process is using them
  - We have to worry about a program's termination without releasing semaphores
  - This can be solved by the semaphore UNDO feature

# Semaphore (3/3)

---

Creating or opening a set of semaphore

- `int semget(key_t key, int nsems, int semflg);`
- Returns: semaphore ID if OK, -1 on error
- Creates a new set of *nsems* semaphores
  - If opening an existing semaphores, this value can be 0
- Upon creating, the least significant 9 bits of *semflg* define the permissions for the semaphore set
- *semflg* can be OR'ed with `IPC_CREAT` and/or `IPC_EXCL`

# Semaphore – System Limitations

---

The limitations may vary on different platforms

- “ipcs -l” command on Linux
- “ipcs -S” on BSD and Mac OS X

```
$ ipcs -l
```

```
...
```

```
----- Semaphore Limits -----
max number of arrays = 128
max semaphores per array = 250
max semaphores system wide = 32000
max ops per semop call = 32
semaphore max value = 32767
```

# Controlling Semaphores (1/3)

---

The internal data structure associated with a semaphore set

```
struct semid_ds {
 struct ipc_perm sem_perm; /* Ownership and permissions */
 time_t sem_otime; /* Last semop time */
 time_t sem_ctime; /* Last change time */
 unsigned short sem_nsems; /* No. of semaphores in set */
};
```

Each member of the semaphore set has at least these attributes maintained by the kernel:

- `semval`: semaphore value, always  $\geq 0$
- `sempid`: pid for last operation
- `semncnt`: # of processes waiting for the `semval` to increase
- `semzcnt`: # of processes waiting for the `semval` to be zero

# Controlling Semaphores (2/3)

---

## Synopsis

- `int semctl(int semid, int semnum, int cmd, /* union semun arg */);`
- Returns: it depends on commands
- This function may be called with 3 or 4 arguments, depends on *cmd*
- The 4<sup>th</sup> argument

```
union semun {
 int val; /* Value for SETVAL */
 struct semid_ds *buf; /* Buffer for IPC_STAT, IPC_SET */
 unsigned short *array; /* Array for GETALL, SETALL */
};
```

# Controlling Semaphores (3/3)

---

## Available *cmds*

| <i>cmds</i> | Description                                                                                                                                 |
|-------------|---------------------------------------------------------------------------------------------------------------------------------------------|
| IPC_STAT    | Retrieve the internal <i>semid_ds</i> data structure and stores in <i>arg.buf</i>                                                           |
| IPC_SET     | Set the internal <i>semid_ds</i> data structure by <i>arg.buf</i><br>■ <i>sem_perm.uid</i> , <i>sem_perm.gid</i> , and <i>sem_perm.mode</i> |
| IPC_RMID    | Remove the semaphore (immediately)                                                                                                          |
| GETVAL      | Return the value of <i>semnum</i> -th member                                                                                                |
| SETVAL      | Set the value of <i>semnum</i> -th member by <i>arg.val</i>                                                                                 |
| GETPID      | Return the value of <i>sempid</i> for the <i>semnum</i> -th member                                                                          |
| GETNCNT     | Return the value of <i>semncnt</i> for the <i>semnum</i> -th member                                                                         |
| GETZCNT     | Return the value of <i>semzcnt</i> for the <i>semnum</i> -th member                                                                         |
| GETALL      | Retrieve all semaphore values, returned by <i>arg.array</i>                                                                                 |
| SETALL      | Set all semaphore values by <i>arg.array</i>                                                                                                |



# Semaphore Operations

---

## Synopsis

- `int semop(int semid, struct sembuf semoparray[], size_t nops);`
- Returns: 0 if OK, -1 on error
- The *semoparray* argument is a pointer to an array of semaphore operations
- Please see the next slide for the details of operations

```
struct sembuf {
 unsigned short sem_num; /* member # in set (0, 1, ..., nsems-1) */
 short sem_op; /* operation (negative, 0, or positive) */
 short sem_flg; /* IPC_NOWAIT, SEM_UNDO */
};
```

# Semaphore Operations – Return Resources

---

*sem\_op* is positive

*sem\_op* is added to the semaphore's value

If SEM\_UNDO is specified, *sem\_op* is *subtracted* from the semaphore's *adjustment value* for this process

# Semaphore Operations – Obtain Resources

---

## *sem\_op* is negative

If resources are available ( $|sem\_op| \leq sem\_val$ )

- $|sem\_op|$  is subtracted from the semaphore's value
- If SEM\_UNDO is specified,  $|sem\_op|$  is *added* to the semaphore's *adjustment value* for this process

If resources are not available ( $|sem\_op| > sem\_val$ )

- If IPC\_NOWAIT is specified, *semop* returns an error of EAGAIN
- If IPC\_NOWAIT is not specified
  - The *semncnt* value for this semaphore is increased
  - The process is suspended until ...
    - The semaphore's value becomes greater than or equal to the  $|sem\_op|$ , the *semncnt* should be increased
    - The semaphore is removed from the system: *semop* returns an error of EIDRM
    - It is interrupted by a signal: *semop* returns an error of EINTR

# Semaphore Operations – Wait until Zero

---

## *sem\_op* is zero

The calling process wants to wait until the semaphore's value becomes 0

If the semaphore's value is currently 0, the function returns immediately

Otherwise,

- If `IPC_NOWAIT` is specified, return is made with an error of `EAGAIN`
- If `IPC_NOWAIT` is not specified
  - The *semzcnt* value for this semaphore is incremented
  - The calling process is suspended until ...
    - The semaphore's value becomes 0 , the *semzcnt* should be increased
    - The semaphore is removed from the system: *semop* returns an error of `EIDRM`
    - It is interrupted by a signal: *semop* returns an error of `EINTR`

# Semaphore Adjustment on Terminating a Process

---

We have mentioned the problem

- A program's termination without releasing semaphores may block future access to the resource

The problem can be solved by the UNDO feature

- When we specify the SEM\_UNDO flag for a semaphore operation
- The kernel remembers how many resources we allocated from that particular semaphore
- When the process terminates, the kernel checks whether the process has any *outstanding semaphore adjustments*, i.e., the value is  $> 0$
- If so, applies the adjustment to the corresponding semaphore
  - *semval* is increased by the adjustments

# Shared Memory

---

Allows two or more processes to share a given region of memory

This is the fastest form of IPC

- The data does not need to be copied between the client and the server, but
- We have to synchronize access to a given region among multiple processes
  - If the server is placing data into a shared memory region, the client should not try to access the data
- Synchronizing can be done by semaphores

# Shared Memory (Cont'd)

---

## Creating or opening a shared memory

### Synopsis

- `int shmget(key_t key, size_t size, int flag);`
- Returns: shared memory ID if OK, -1 on error
- Upon creating, the least significant 9 bits of *semflg* define the permissions for the shared memory
- *flag* can be OR'ed with `IPC_CREAT` and/or `IPC_EXCL`
- The actual size of the created shared memory is round up to multiples of the `PAGE_SIZE` (4096 bytes)
- When a shared memory is created, it's content initialized to all zero

# Shared Memory – System Limitations

---

The limitations may vary on different platforms

- “ipcs -l” command on Linux
- “ipcs -M” on BSD and Mac OS X

```
$ ipcs -l
```

```
...
```

```
----- Shared Memory Limits -----
```

```
max number of segments = 4096
```

```
max seg size (kbytes) = 18014398509465599
```

```
max total shared memory (kbytes) = 18446744073642442748
```

```
min seg size (bytes) = 1
```



# Controlling Shared Memory

---

The internal data structure associated with a shared memory

```
struct shmid_ds {
 struct ipc_perm shm_perm; /* Ownership and permissions */
 size_t shm_segsz; /* Size of segment (bytes) */
 time_t shm_atime; /* Last attach time */
 time_t shm_dtime; /* Last detach time */
 time_t shm_ctime; /* Last change time */
 pid_t shm_cpid; /* PID of creator */
 pid_t shm_lpid; /* PID of last shmat(2)/shmdt(2) */
 shmatt_t shm_nattch; /* No. of current attaches */
 ...
};
```

# Controlling Shared Memory (Cont'd)

---

## Synopsis

- `int shmctl(int shmid, int cmd, struct shmid_ds *buf);`
- Returns: 0 if OK, -1 on error
- Commands
  - IPC\_STAT: Retrieve the internal `shmid_ds` data structure
  - IPC\_SET: Set the internal `shmid_ds` data structure
    - `shm_perm.uid`, `shm_perm.gid`, and `shm_perm.mode`
  - IPC\_RMID: Remove the shared memory, but *it is **actually removed** until the last process using the segment terminates or detaches it*
  - SHM\_LOCK: Make the shared memory not swappable
  - SHM\_UNLOCK: Make the shared memory swappable
    - The last two commands can be only used by superuser

# Attach a Shared Memory

---

## Synopsis

- `void *shmat(int shmid, const void *addr, int flag);`
- Returns: pointer to shared memory segment if OK, -1 on error
- The *addr* argument
  - If *addr* is NULL, the segment is attached at the first available address selected by the kernel (*\*RECOMMENDED*)
  - If *addr* is not NULL and SHM\_RND is not specified, the segment is attached at the address given by *addr*
  - If *addr* is not NULL and SHM\_RND is specified, the segment is attached at the address given by  $(addr - (addr \bmod SHMLBA))$ 
    - Round down to the multiples of SHMLBA
- The *flag* argument
  - If the SHM\_RDONLY bit is specified in *flag*, the segment is attached read-only

# Detach a Shared Memory

---

## Synopsis

- `int shmdt(void *addr);`
- Returns: 0 if OK, -1 on error

# Q & A

---